Streaming-Order AMM

FRE-GY 6931 Blockchain & Cryptocurrency, Fall 2026 · Topic 1: Automated Market Maker · Final project report

Team: Aditya Shah (as22008), Siddhant Yadav (sy4583), Dev Salvi (ds8259)

Network: Ethereum Sepolia · Repository: https://github.com/M720d/Streaming-Order-AMM (tag `final`) · Video: linked in `submission.json`

Contents

1. Objective
   1.1 Why we chose this, and what we expected
   1.2 What is original and what is not
2. Design decisions
   2.1 One price unit across the system (changed during design)
   2.2 Slicing: fixed count, equal slices, no dust
   2.3 Oracle guard: skip when the pool is off, revert when the oracle is bad
   2.4 Slippage: an oracle-derived minimum, per order
   2.5 The order is an ERC-721 and its holder owns everything
   2.6 Trust and admin surface kept as small as possible
   2.7 Gas: one change that worked, one that did not
   2.8 What we did not build, and one deliberate difference from the proposal
   2.9 What changed because we ran it on Sepolia
3. Architecture
   3.1 Contracts and how they connect
   3.2 The pool-side contracts
   3.3 The order-side contracts
   3.4 Off-chain components
   3.5 Address table (Sepolia, all verified on Etherscan)
4. Financial analysis
   4.1 Price impact by trade size (Part A)
   4.2 Lump sum versus 5, 10 and 20 slices (Part B)
   4.3 Interaction with the oracle guard
   4.4 Impermanent loss and LP profit (Part A)
   4.5 Reconciliation: model versus contracts
   4.6 What the model leaves out
   4.7 Cost per user action and break-even (Part C)
5. Testing, security and gas
   5.1 Tests
   5.2 Slither
   5.3 Gas
   5.4 What the testing does not cover
6. Going to market
   6.1 Users
   6.2 Trust
   6.3 The first hundred users
   6.4 Business model with numbers
   6.5 Risks
7. Summary
8. AI Use
Appendix A. Execution evidence: Etherscan links and screenshots
   A.1 Transactions
   A.2 Keeper log (order 2, the order created in the web app)
   A.3 Screenshots
Appendix B. Source code, metadata and JSON
References

1. Objective

Selling a large amount of a token into a constant-product pool is expensive: the price moves against the seller with every unit sold, so the last coins sell far below the first. Selling 20 WBTC into a pool of 500 WBTC and 30 M USDC in one swap returns 1,150,517 USDC, 4.1% below what the same coins are worth at the pool's starting price, and most of that is price impact, not the 0.3% fee. The usual answer is to split the order. Splitting only helps, though, if someone puts the price back between the pieces; otherwise each piece sells into the price the previous one left behind, and the total is the same as one big trade.

We built a small exchange that does the splitting for the seller and measured when it actually helps:

- a constant-product **Pool** for mock WBTC and USDC with a 30 bps fee, LP shares and a time-weighted average price;
- an **OrderBook** where a seller deposits the whole order once and receives an ERC-721 that owns it; the order is sold in a fixed number of equal slices, one per interval, by anyone who calls `executeSlice` (our keeper script does);
- a **PriceGuard** that reads Chainlink BTC/USD, refuses to trade on a stale, zero or carried-over price, and skips a slice when the pool has drifted too far from Chainlink;

- off-chain, a **keeper**, an **arbitrage bot** that refills the pool between slices, and a **web app** to create, watch, claim and cancel orders.

The research question is the one in the proposal: **how much does slicing save compared with a lump sum, and when does the saving disappear?** Section 4 answers it with notebooks whose integer model reproduces the contracts exactly, and Part C adds what the saving is worth after gas and whether a business could run on it. Everything runs on Ethereum's Sepolia testnet with mock tokens (Section 3.5 lists the verified contracts); no real money is involved.

1.1 Why we chose this, and what we expected

We took Topic 1 (automated market maker) because it let us test a claim from our own proposal with numbers instead of intuition: that splitting a large sale into slices only helps if the market refills the pool between slices. Before building we expected three things: (1) without a refill, slicing saves nothing; (2) with a refill it saves a lot, with diminishing returns as slices increase; and (3) that the practical limits would be the oracle check and gas, not the maths. The results in Sections 4 and 5 confirm the first two and add a surprise to the third (Sepolia gas is higher than our local tests).

1.2 What is original and what is not

The idea of executing a large order over time is not new: CoW Swap's TWAP orders, Uniswap X and Paradigm's TWAMM address the same problem. What is ours is the combination and the evidence: the order is an **ERC-721 position** whose holder owns the unsold balance and the proceeds and can transfer it; the guard treats a bad market (skip) differently from a bad oracle (revert); the analysis uses the contract's own integer arithmetic, so the model matches the contracts exactly, locally (56 comparisons) and on Sepolia (3 of 3 slices); and the business case states plainly that the fees we proposed are not built and what that costs.

2. Design decisions

The system is four contracts on top of two test tokens: a constant-product `Pool` (also the LP share token), a `PriceGuard` that reads Chainlink BTC/USD, and an `OrderBook` that turns a large sell into an ERC-721 and executes it in equal slices. This section records the choices that shaped the contracts, what we considered, and what we changed after finding a problem.

2.1 One price unit across the system (changed during design)

First idea. Let the pool report its raw reserve ratio (raw USDC per raw WBTC, scaled by 1e18) and have the guard compare it with the oracle.

Problem. With WBTC at 8 decimals and USDC at 6, that ratio is 100 times smaller than the oracle's USD-per-BTC price (a pool at 60,000 USDC/BTC reports 6e17 where the oracle says 6e22). A guard comparing the two would reject every pool price, or accept a pool that is 100 times wrong, depending on which side was scaled.

Decision. Every price in the system is **USDC per one whole WBTC, scaled by 1e18** (60,000 USDC/BTC is `60_000e18`). The pool converts internally and raw-unit ratios never leave it. All remaining conversions on the order side (Chainlink's 8-decimal answer into this unit, and a WBTC amount into the raw USDC it is worth) live in one library, `PriceMath`, with its own tests, so there is a single place to audit. Interface functions carry the unit in their names (`spotPriceE18`, `consultTwapE18`, `poolWithinBounds(poolPriceE18)`).

2.2 Slicing: fixed count, equal slices, no dust

The slice count is a fixed integer (1 to 200) chosen at creation; there is no continuous TWAMM-style streaming. Each slice sells `total / slices`, and the **last slice sells everything that is left**, so rounding remainders are never stranded in the contract (tested with 10 WBTC + 3 raw units in 5 slices). The alternative, keeping dust and letting the seller sweep it, adds a state and a way to lose funds for no benefit.

Timing. The first slice is due one interval after creation. After every attempt, successful or skipped, the next one is due one interval **from now**, not one interval after the previous due time. We considered scheduling on a fixed grid,

but a keeper that is late would then fire several slices back to back, which removes the refill time that slicing depends on (Section 4).

2.3 Oracle guard: skip when the pool is off, revert when the oracle is bad

Two different failures get two different behaviours.

- **Pool price outside the allowed band around the oracle** (default 5%, hard cap 20%): the slice is **skipped** with a `SliceSkipped` event, no tokens move, and the order tries again an interval later. Nothing is wrong with the order; the market is just temporarily off, and a revert would only waste the keeper's gas and hide the reason.
- **Oracle unusable** (price zero or negative, older than the staleness limit or dated in the future, or `answeredInRound` behind `roundId`): the call **reverts**. We cannot tell whether the pool is off, so trading blind is the one thing the guard must never do. This is the proposal's "no trading when the feed is stale".

The guard compares the pool's **spot** price with the oracle. We considered using the pool's own TWAP for the pool side, which resists single-block manipulation, but a TWAP lags and needs history that a young pool does not have. Spot against a fresh oracle is a tighter check at the moment of the trade, and the residual risk is covered by the minimum-output rule below.

2.4 Slippage: an oracle-derived minimum, per order

Each slice must return at least the oracle's quote for that slice minus the order's `toleranceBps` (capped at 10%), otherwise the pool call reverts. The seller chooses the tolerance. We learned two things from testing:

1. **Tolerance must cover the fee plus the accumulated drift.** An early test used 3% on a 10 WBTC order in five slices; the last slice reverted because, without a refill, it returned 116,365 USDC against a minimum of 116,400 (a 3.03% shortfall from the oracle quote of 120,000 for its 2 WBTC). We raised the test tolerance to 5% and wrote down the rule in the contract's NatSpec: tolerance has to cover the 30 bps fee plus the expected impact, or the slice reverts.
2. **Skip and revert coexist.** A slice that is too bad for the seller's tolerance reverts; it does not skip. The seller can always cancel.

Known limitation. A searcher can sandwich a slice. The attacker can push the pool price up to the guard band (5% by default) and the seller still receives at least the oracle quote minus tolerance, so the worst-case loss per slice is bounded by the tolerance and band the seller and owner chose, not unbounded. We did not add a private-mempool or commit-reveal scheme.

2.5 The order is an ERC-721 and its holder owns everything

The NFT holder owns the unsold WBTC and the USDC proceeds. `claim` pays the **current** `ownerOf`, so a transferred order pays its new holder (tested: slices run, the NFT moves, the new holder claims and the old one cannot cancel). Further choices:

- `claim` and `executeSlice` are callable by anyone, because the money can only go to the holder and a permissionless keeper is part of the design. `cancel` is **holder only**, not an approved operator, because it releases all funds and burns the NFT.
- `cancel` returns the unsold WBTC *and* any unclaimed USDC, deletes the order and burns the NFT. On a finished order it works as claim-and-burn.
- Minting uses `_safeMint` so contract wallets (multisigs, treasuries) can receive orders. That introduces a receiver callback, which is why every state-changing function is `nonReentrant` (tested by a receiver that tries to cancel during the mint, and a pool that calls back during the swap).

2.6 Trust and admin surface kept as small as possible

- `OrderBook` and `Pool` have **no owner and no admin function**. The pool and guard addresses are immutable.
- `PriceGuard` has one owner who can change exactly two numbers: the staleness limit (1 second to 1 day) and the deviation band (1 to 2,000 bps). Both have hard caps, emit events, and zero is rejected so a mistake

cannot brick the guard. The owner cannot change the feed or touch funds. The same list is in the contract's NatSpec.
- The `OrderBook` approves the pool for exactly the slice amount before each swap. We considered a one-time unlimited approval, which saves an estimated 20k gas per slice (not measured), and rejected it: it would expose every order's WBTC to the pool address if the pool ever had a bug.

2.7 Gas: one change that worked, one that did not

`createOrder`, `executeSlice`, `claim` and `cancel` were measured before and after (`docs/gas-orders.md`). The order struct in the interface needs three storage slots only because the 16-bit tolerance spills over. Keeping the interface unchanged but storing the order in two slots (narrowing the slice size to 96 bits, about 1e13 times the WBTC supply, and rejecting anything larger) saves **22,054 gas on `createOrder`** (11.9%), 1,986 on each slice and 4,005 on `cancel`. Our first attempt, caching the order in memory inside `executeSlice`, made it **924 gas worse** and was reverted.

2.8 What we did not build, and one deliberate difference from the proposal

- **Execution fee and protocol share of the swap fee.** The proposal's revenue model has both. We did not build them: the contracts charge nothing beyond the pool's 30 bps, which goes entirely to LPs, so **as built the protocol earns nothing**. We treated both as future work rather than change tested contracts days before the deadline. Sections 4.7, 6.4 and 7 say what that means for the business.
- **Early-LP rewards paid in the share token** (the proposal's plan for the first hundred users) are not built; that needs a reward contract (Section 6.3).
- **The LP token is the `Pool` contract itself**, not a separate ERC-20 as the proposal's contract table listed. That is one contract fewer to deploy and review, and share accounting and reserves change in the same call. The cost is that the pool's address is also the token's address, so the pool cannot be replaced while keeping the same token. Nothing in the tests depends on the choice.
- The keeper, the arbitrage bot and the front end are described in Section 3.4.

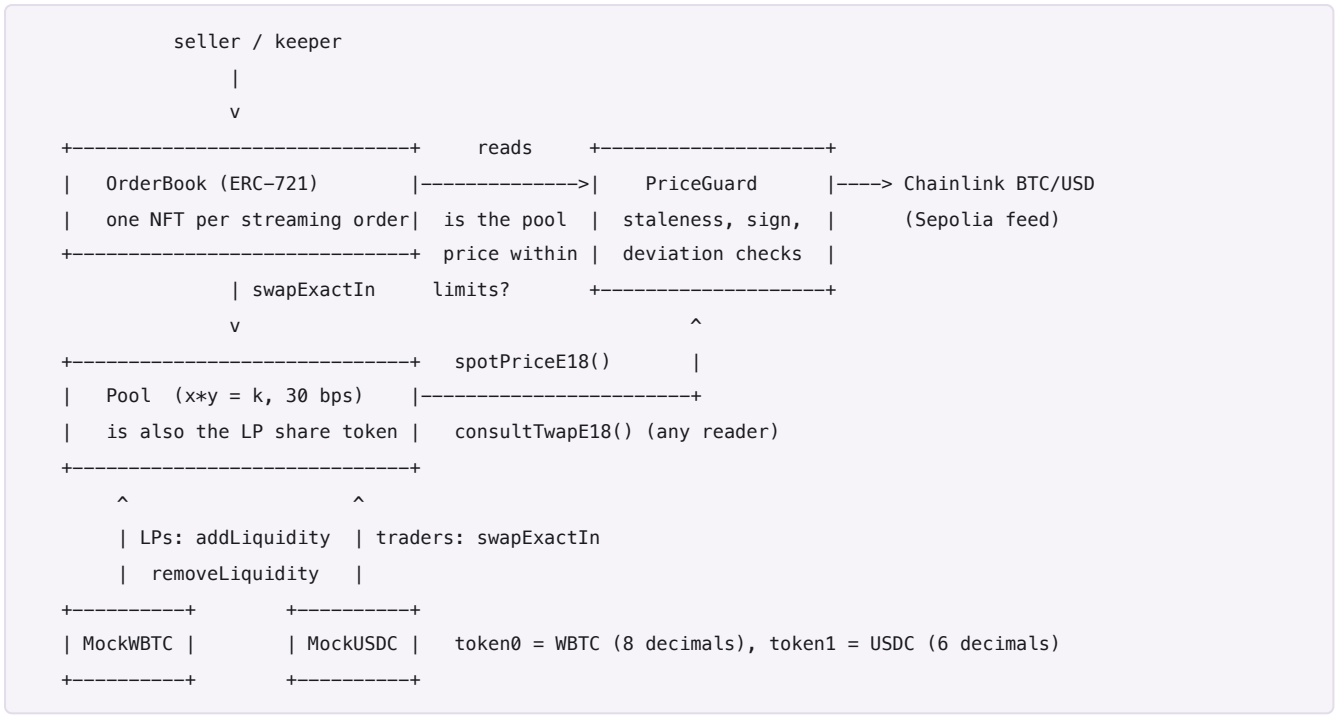
2.9 What changed because we ran it on Sepolia

- **Staleness limit 3,600 s became 7,200 s.** Chainlink's Sepolia BTC/USD feed updated at 20:44 and 21:44 UTC on 2026-11-01, so it updates hourly and its age cycles between zero and one hour. A one-hour limit would make slices revert near the end of each cycle. The owner can raise the limit up to the hard cap of one day.
- **Arbitrage threshold 50 bps became 35 bps.** A first live order (1 WBTC in three slices) moved the pool only 39 bps from Chainlink, so a 50 bps bot never traded. The pool fee is 30 bps, so 35 is about the smallest threshold that can still pay.
- **Front-end reads.** Alchemy's free tier allows `eth_getLogs` over only 10 blocks, so order discovery starts at the deploy block and Sepolia reads go through a public RPC with Alchemy as the fallback. The Chainlink feed address is read from `PriceGuard.feed()` instead of looking up the local-only mock aggregator.
- **Gas.** Actions that create storage cost 2 to 4 times more gas on Sepolia than in our local tests, and contract creation about 7 times more (Section 5.3). Our first cost estimates assumed they would match.

3. Architecture

3.1 Contracts and how they connect

The system is six contracts on Sepolia: two test tokens, a pool, an oracle guard, an order book, and Chainlink's BTC/USD feed, which we read but do not own.



The pool is the only contract that holds liquidity and the only one that sets prices. The order book never computes a price itself; it asks the pool for a swap and asks the guard whether the pool's price is acceptable. Section 3.3 covers the order side.

3.2 The pool-side contracts

MockWBTC and MockUSDC are test tokens (OpenZeppelin ERC-20) with 8 and 6 decimals, matching the real assets. Sepolia has a BTC/USD price feed but no liquid WBTC/USDC market, so we issue our own pair. Each has an owner-only `mint` and a public `faucet(amount)` capped per call (5 WBTC, 500,000 USDC) so teammates can fund themselves without asking the owner. These are the only owner powers in the pool-side contracts, and they exist for testing only.

Pool is a constant-product pool for the pair, and it is also the ERC-20 that represents LP shares (symbol `SALP`). Everything a liquidity provider or trader does goes through five functions:

Function	What it does
`addLiquidity(a0Desired, a1Desired, a0Min, a1Min, to)`	Takes tokens in the pool's current ratio (never less than the minimums) and mints shares to `to`. The first deposit mints `sqrt(a0 * a1) - 1000` shares and locks 1000 forever, so the pool can never be emptied and the price stays defined.
`removeLiquidity(shares, a0Min, a1Min, to)`	Burns the caller's shares and sends the matching share of both reserves to `to`.
`swapExactIn(tokenIn, amountIn, minAmountOut, to)`	Sells an exact amount in, reverts with `SlippageExceeded` if the output is below `minAmountOut`. The 30 bps fee is deducted from the input and stays in the reserves, so it accrues to LPs and `k` grows on every trade.
`spotPriceE18()`	Current price, USDC per whole WBTC, scaled by 1e18 (60,000 USDC/BTC reads as `60_000e18`).
`consultTwapE18(window)`	Time-weighted average of that price over the last `window` seconds.

Design points specific to the pool:

- **Integer maths only, floor division.** The swap formula is `out = in*9970 * reserveOut / (reserveIn*10000 + in*9970)`. Because there is no floating point, the Python model in `docs/analysis_pool.ipynb` reproduces every contract output exactly, and the Hardhat tests assert values computed in Python rather than values read back from the contract.
- **TWAP from the pool's own accumulator.** On every reserve change the pool adds `price * seconds elapsed` to a running total and records `(timestamp, total)`. `consultTwapE18` finds the record at the start of

the window, interpolates to the exact start using the price that was in force then, and divides the difference by the window. Records are stored once per block timestamp, packed into one 256-bit slot. A window longer than the pool's age reverts with `InsufficientHistory` instead of returning a biased number.

- **No admin.** The pool has no owner, no pause, no upgrade path and no fee switch. The fee is a constant. Nothing in it can be changed after deployment.
- **One price unit.** The pool reads both tokens' decimals at construction and converts internally, so the raw reserve ratio never leaves the contract (Section 2.1 explains why this mattered).
- **Reserves and update time share one storage slot** (112 + 112 + 32 bits). This is the gas optimization in Section 5.3; it caps each reserve at about 5.2e33 raw units, enforced with a checked cast.

What the pool does not support: fee-on-transfer and rebasing tokens (it credits the amount it asked for), and ERC-777-style token callbacks. Both tokens here are our own plain ERC-20s.

3.3 The order-side contracts

OrderBook is an ERC-721 (`SORDER`) with no owner. Every order is one NFT; its holder owns the unsold WBTC and the USDC proceeds.

Function	What it does
`createOrder(totalIn, slices, interval, toleranceBps)`	Pulls `totalIn` WBTC from the caller (needs an approval), stores the order and mints the NFT. Slices 1 to 200, tolerance at most 10%.
`executeSlice(orderId)`	Callable by anyone once the interval has elapsed. Asks PriceGuard whether the pool is inside the band; if not, emits `SliceSkipped` and retries one interval later. Otherwise sells `total / slices` (the last slice sells everything left) through `Pool.swapExactIn`, with a minimum output of the oracle's quote minus the order's tolerance.
`claim(orderId)`	Pays the accumulated USDC to the current NFT holder.
`cancel(orderId)`	Holder only. Returns the unsold WBTC and any unclaimed USDC and burns the NFT.
`getOrder(orderId)`	View: remaining amount, proceeds, slice size, slices left, interval, next execution time, tolerance.

Every state-changing function is `nonReentrant`, all transfers use SafeERC20, and the pool is approved for exactly the slice amount before each swap.

PriceGuard reads Chainlink and answers one question for the order book.

Function	What it does
`getOraclePrice()`	Returns the BTC/USD price in the system's price unit (USDC per WBTC, 1e18). Reverts if the answer is zero or negative, older than `maxStaleness`, dated in the future, or from a carried-over round (`answeredInRound < roundId`).
`poolWithinBounds(poolPriceE18)`	True when the pool price is within `maxDeviationBps` of the oracle price.
`setMaxStaleness`, `setMaxDeviationBps`	**The only admin functions in the whole system.** Owner only, hard-capped (1 s to 1 day; 1 to 2,000 bps), emit events. The owner cannot change the feed or touch funds.

3.4 Off-chain components

- **Keeper** (`packages/hardhat/scripts/keeper.ts`). Finds orders from `OrderCreated` events and, for each order whose interval has elapsed, simulates `executeSlice` first and sends it only if the simulation would execute. It sends one transaction at a time, never re-sends a transaction that is still pending, and logs `filled`, `skipped` or `reverted` with the transaction hash.
- **Arbitrage bot** (`scripts/arb.ts`). Reads the Chainlink price and the pool price. When the gap exceeds a threshold it sizes the profit-maximising trade (the same rule as the analysis notebook), caps it by a maximum size and the wallet balance, and trades only if the expected profit after the 30 bps fee is above a minimum. It approves exactly the amount it trades. This is what refills the pool between slices.

- **Front end** (`packages/nextjs`, Scaffold-ETH 2). Pool and oracle view, a create-order form with a lump-sum versus sliced preview and a confirmation step that lists what the user will sign, order progress, claim and cancel, transaction status, plain language error messages, and a warning when the Chainlink feed is stale.
- **Tooling.** `yarn fund-arb`, `yarn send-eth` and `yarn evidence` (a read-only script that lists deploy and first-event transactions with gas); the run-book is `docs/sepolia-runbook.md`.

3.5 Address table (Sepolia, all verified on Etherscan)

Contract	Role	Address (source verified)	Deployment tx
MockWBTC	Test token, 8 decimals, capped faucet	`0x2c95a78b3a3fc73c16bdd6394bfe30cc51d37d12`	`0x688d1fd6...`
MockUSDC	Test token, 6 decimals, capped faucet	`0x9cd5587b22d15790a5880d37837cdd091c9ea0b5`	`0xae7c03dc...`
Pool	Constant-product pool, 30 bps, TWAP; **also the ERC-20 LP token `SALP`**	`0x9461fd23335242663ddf7a9fae994995df7fc8e9`	`0xa9f51895...`
PriceGuard	Chainlink checks and pool deviation band	`0xe52d276bec7b96bcdd0632050cad5717956289df`	`0xe52450a3...`
OrderBook	**ERC-721 `SORDER`**, streaming orders	`0x9c20dbb0bb17a0532348e69502af3b8b44fefa62`	`0xbb0acbd1...`
Chainlink BTC/USD feed (external)	Price oracle, 8 decimals	`0x1b44F3514812d835EB1BDB0acB33d3fA3351Ee43`	n/a

The tokens we mint are the **SALP** ERC-20 (LP shares, minted by `Pool.addLiquidity`) and the **SORDER** ERC-721 (one per order, minted by `OrderBook.createOrder`); both are used by the application. Seed liquidity: `0xf2bea583...`.

4. Financial analysis

Two notebooks answer one question: how much does slicing a large order save compared with selling it at once, and when does the saving disappear? `docs/analysis_pool.ipynb` (Part A) covers the pool; `docs/analysis_slicing.ipynb` (Part B) covers slicing. Both use the contract's integer arithmetic (30 bps fee taken from the input and left in the reserves, floor division), so each number below is what the contracts compute. Starting pool: 500 WBTC and 30,000,000 USDC (60,000 USDC per WBTC), as in the proposal.

4.1 Price impact by trade size (Part A)

Trade (WBTC)	% of reserve	Impact, 30 bps fee	Impact, no fee
0.05	0.01%	0.31%	0.01%
2.5	0.5%	0.79%	0.50%
12.5	2.5%	2.72%	2.44%
50	10%	9.34%	9.09%
100	20%	16.88%	16.67%

Impact is driven by the trade as a share of the reserve. The proposal's figures (about 9% for 50 BTC, about 0.5% for a 2.5 BTC slice) are the no-fee column; with the 30 bps fee they are 9.3% and 0.8%. A 50 BTC sale costs 9.3% in a 500 BTC pool and 1.3% in a 5,000 BTC pool.

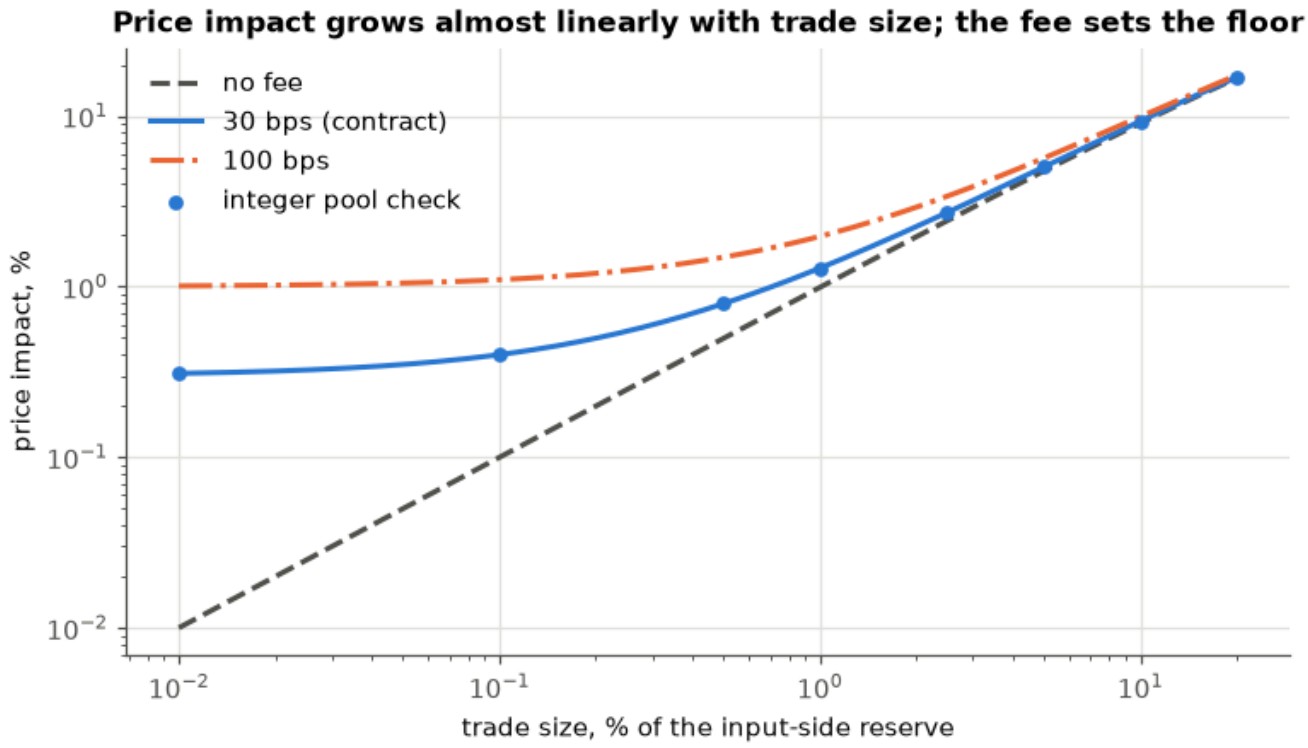


Figure 1. Price impact grows almost linearly with trade size; the fee sets the floor.

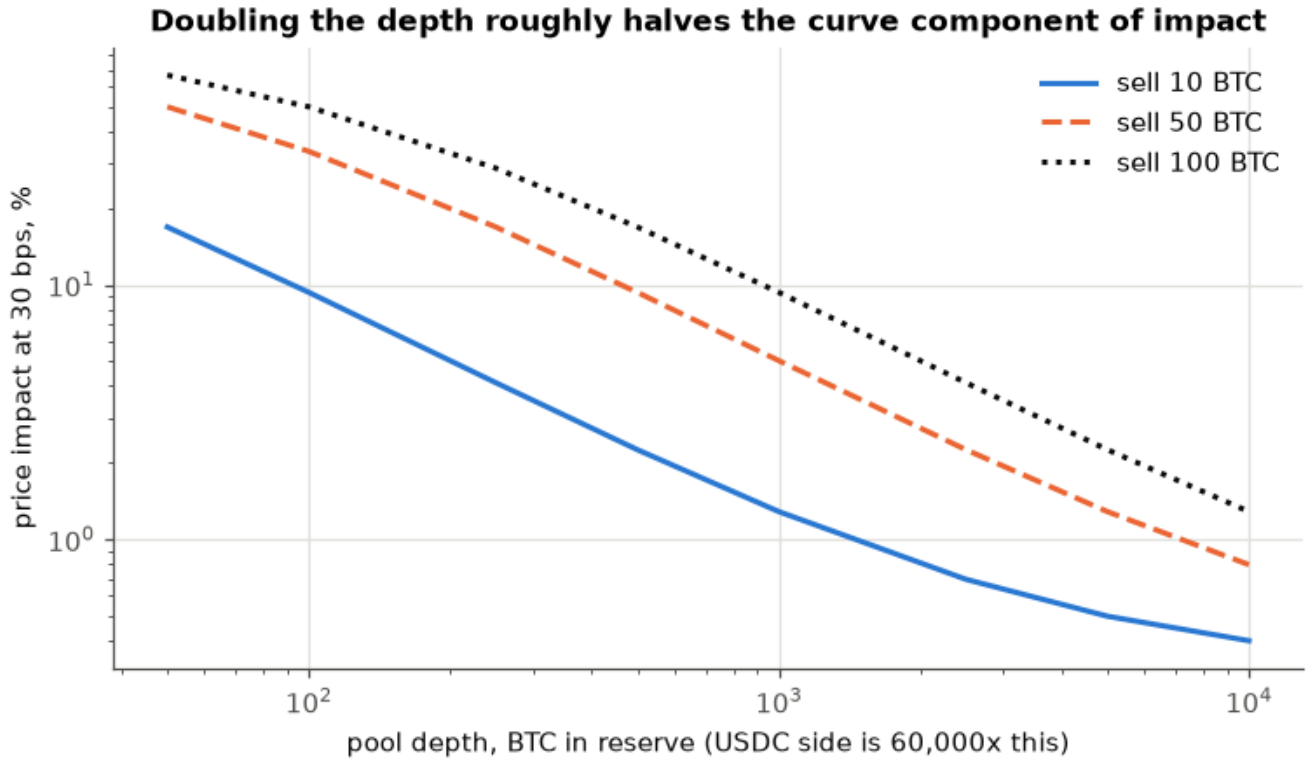


Figure 2. Doubling the depth roughly halves the curve component of impact.

4.2 Lump sum versus 5, 10 and 20 slices (Part B)

Selling 50 WBTC (10% of the pool), cost measured against the frictionless value 50 x 60,000 USDC:

Strategy	Proceeds (USDC)	Cost vs frictionless	Saving vs lump sum
Lump sum	2,719,832.68	934 bps	n/a
5 slices, no refill	2,719,547.22	935 bps	-1.05 bps

Strategy	Proceeds (USDC)	Cost vs frictionless	Saving vs lump sum
10 slices, no refill	2,719,510.48	935 bps	-1.18 bps
20 slices, no refill	2,719,492.02	935 bps	-1.25 bps
5 slices, refill	2,925,696.22	248 bps	+757 bps
10 slices, refill	2,953,598.46	155 bps	+859 bps
20 slices, refill	2,967,746.04	108 bps	+912 bps

"Refill" means an arbitrageur restores the pool to its no-profit price at the external price between slices, at no cost to the seller.

Without a refill, slicing does not help. The pool depends only on how much WBTC is in the reserve, not on how it arrived, so the price ends in the same place. With zero fee the sliced and lump outputs differ by a few raw units of rounding (-1, -3 and -9 raw USDC for 5, 10 and 20 slices). With the 30 bps fee the slices are 1.05 to 1.25 bps *worse*, because each slice's fee stays in the reserve and worsens the next slice's price.

With a refill, slicing saves a lot, with diminishing returns: 757, 859 and 912 bps for 5, 10 and 20 slices (about 9% of proceeds at 20).

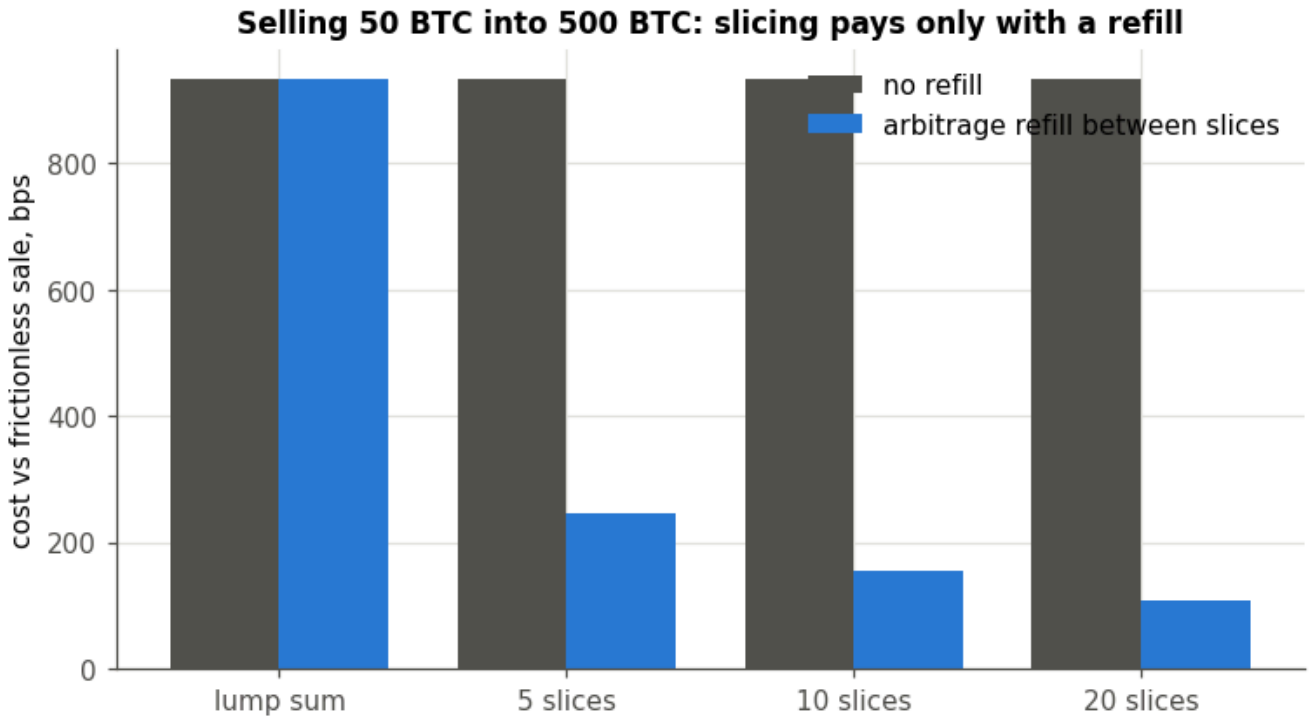


Figure 3. Selling 50 BTC into 500 BTC: slicing pays only with a refill.

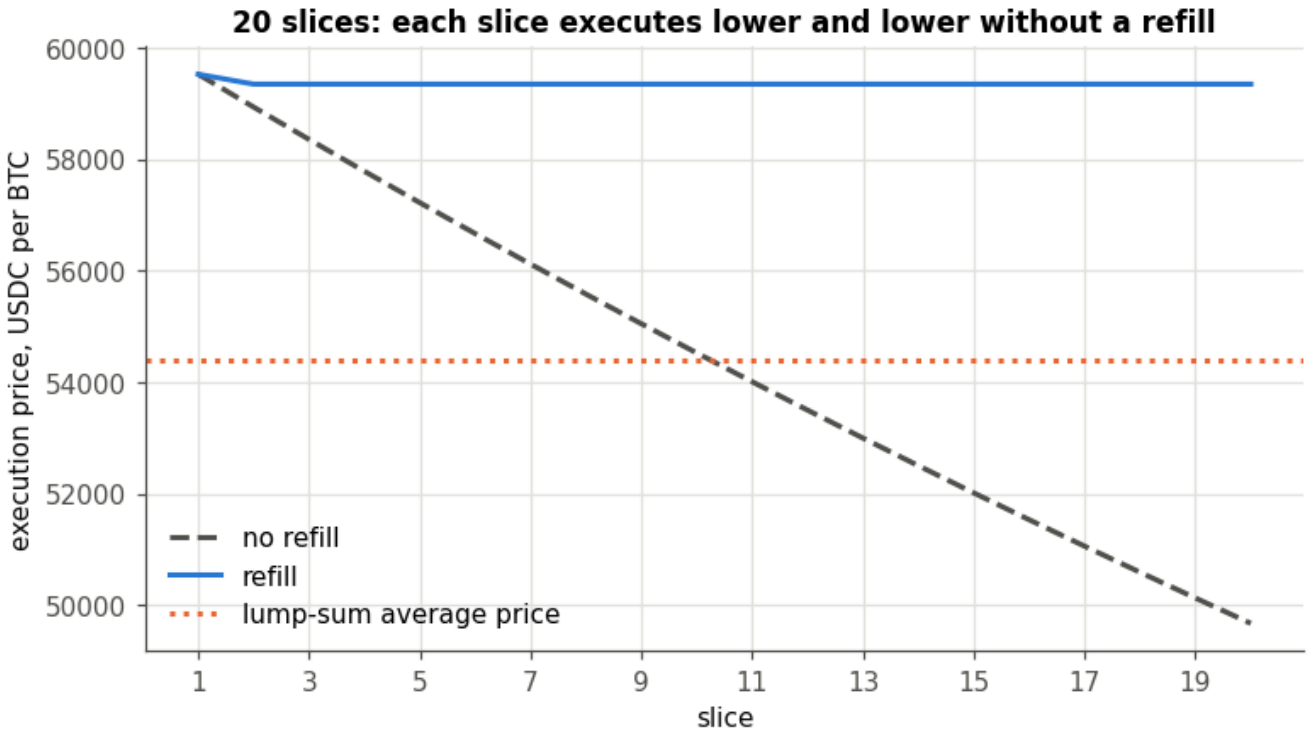


Figure 4. 20 slices: each slice executes lower and lower without a refill.

The saving does not vanish gradually. Letting `rho` be the share of the price gap that is closed between slices:

Gap closed per slice	Saving, 5 slices	Saving, 20 slices
0.1%	0.4 bps	4.6 bps
1%	13 bps	56 bps
10%	133 bps	407 bps
50%	514 bps	817 bps
100%	757 bps	912 bps

The curve is steep at first and then flattens: 10% of the refill already gives 45% of the 20-slice saving. The saving disappears only when almost nothing refills the pool between slices. The slice interval matters through this share, because a longer interval gives arbitrageurs time to act. The notebook uses the refill share as a stand-in for the interval; the Sepolia order in Section 4.5 shows one real case, in which the arbitrage bot, polled about every 20 seconds, restored the pool to within the fee band of Chainlink before every slice.

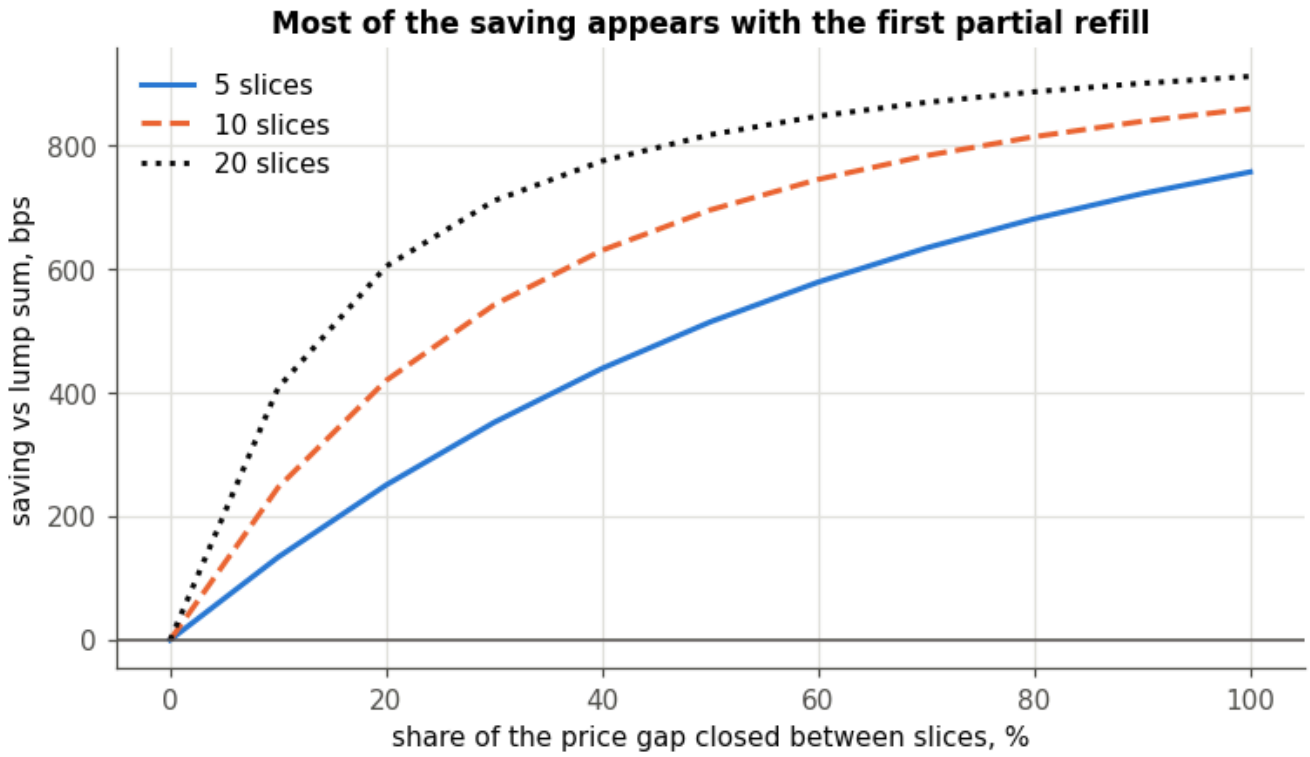


Figure 5. Most of the saving appears with the first partial refill.

Depth and fee (20 slices, full refill, saving vs lump sum): 4,581 bps in a 100 BTC pool, 912 bps at 500 BTC, 161 bps at 2,500 BTC; at 500 BTC the fee changes it only from 945 bps (no fee) to 833 bps (100 bps fee). Slicing saves curve impact, not fees, so the fee barely matters, but pool depth decides how much there is to save.

4.3 Interaction with the oracle guard

Without a refill, the 5% guard first skips at **slice 7 of 20**, and slice 20 would need 1,720 bps of tolerance, above the 1,000 bps cap. So a 50 BTC order in 20 slices **cannot finish without a refill, whatever tolerance the seller sets**. This is the guard working as intended, and it is why the demo uses our own arbitrage script to refill the thin Sepolia pool.

4.4 Impermanent loss and LP profit (Part A)

Impermanent loss depends only on the price move `r`: 5.72% at 2x or 0.5x, 20.0% at 4x, 42.5% at 10x, matching the closed form `2*sqrt(r)/(1+r) - 1` exactly in the integer pool. The 30 bps fee opens a no-trade band of about +/-0.3% around the start price and otherwise trims the loss by only 0.04 to 0.06 points (5.72% becomes 5.68% at 2x). Over a simulated 30-day path in the base case, fees average 1.07% of TVL against a mean IL of 0.38%, so the LP ends on average 0.68% ahead of simply holding and beats holding on 91.6% of paths. The break-even organic volume that covers expected IL is about 0.4%, 2.4% and 8.1% of TVL per day at 30%, 60% and 100% volatility. This is Part A's result; see the notebook for assumptions (price path, volume model, arbitrage timing).

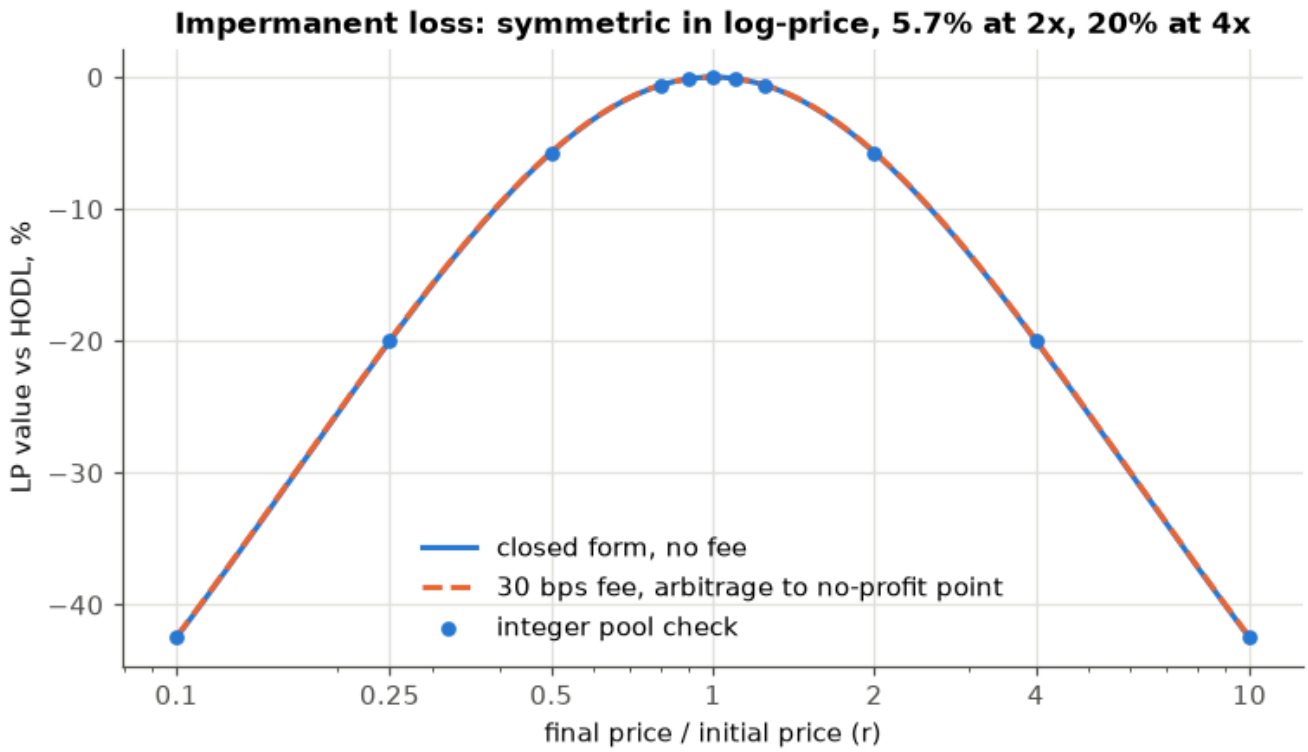


Figure 6. Impermanent loss is symmetric in log-price: 5.7% at 2x, 20% at 4x.

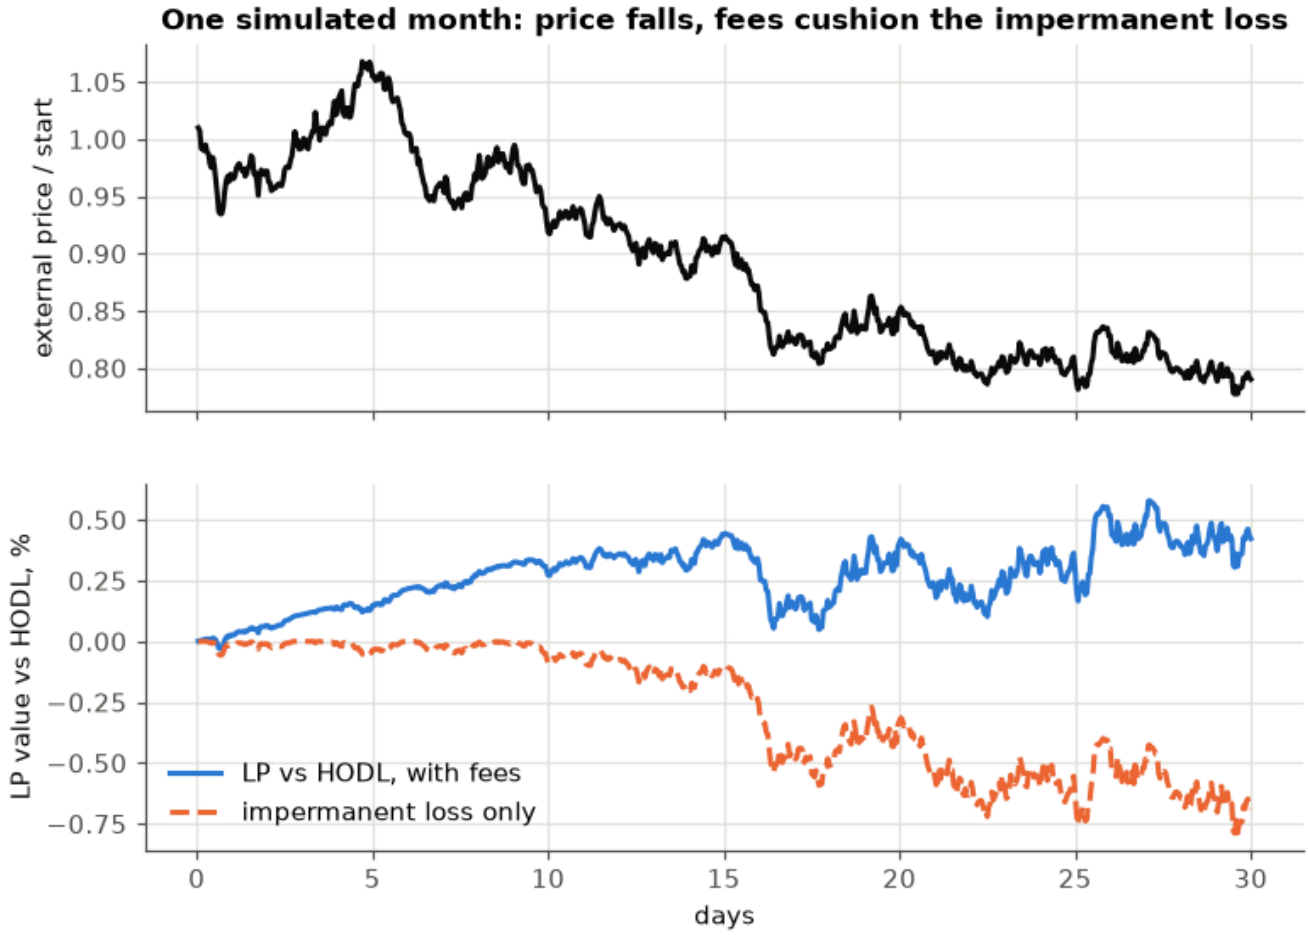


Figure 7. One simulated month: price falls, fees cushion the impermanent loss.

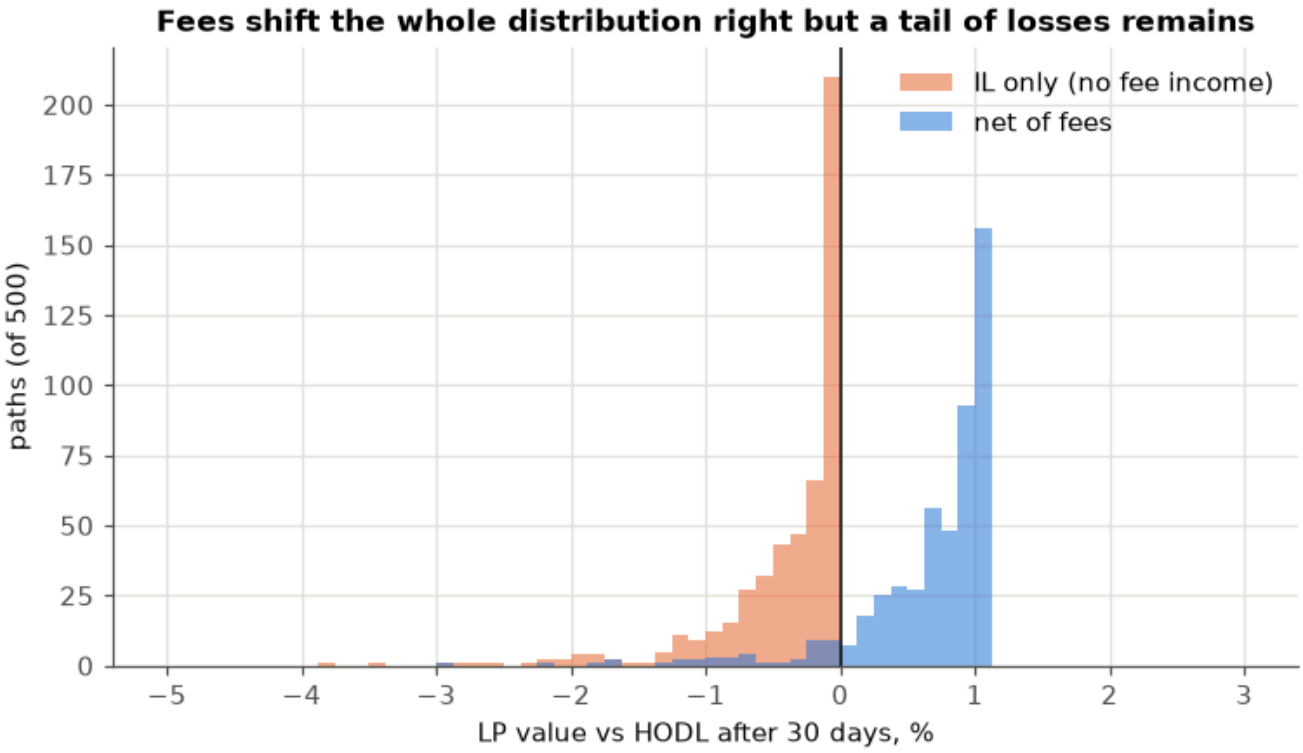


Figure 8. Fees shift the whole distribution of LP results right, but a tail of losses remains.

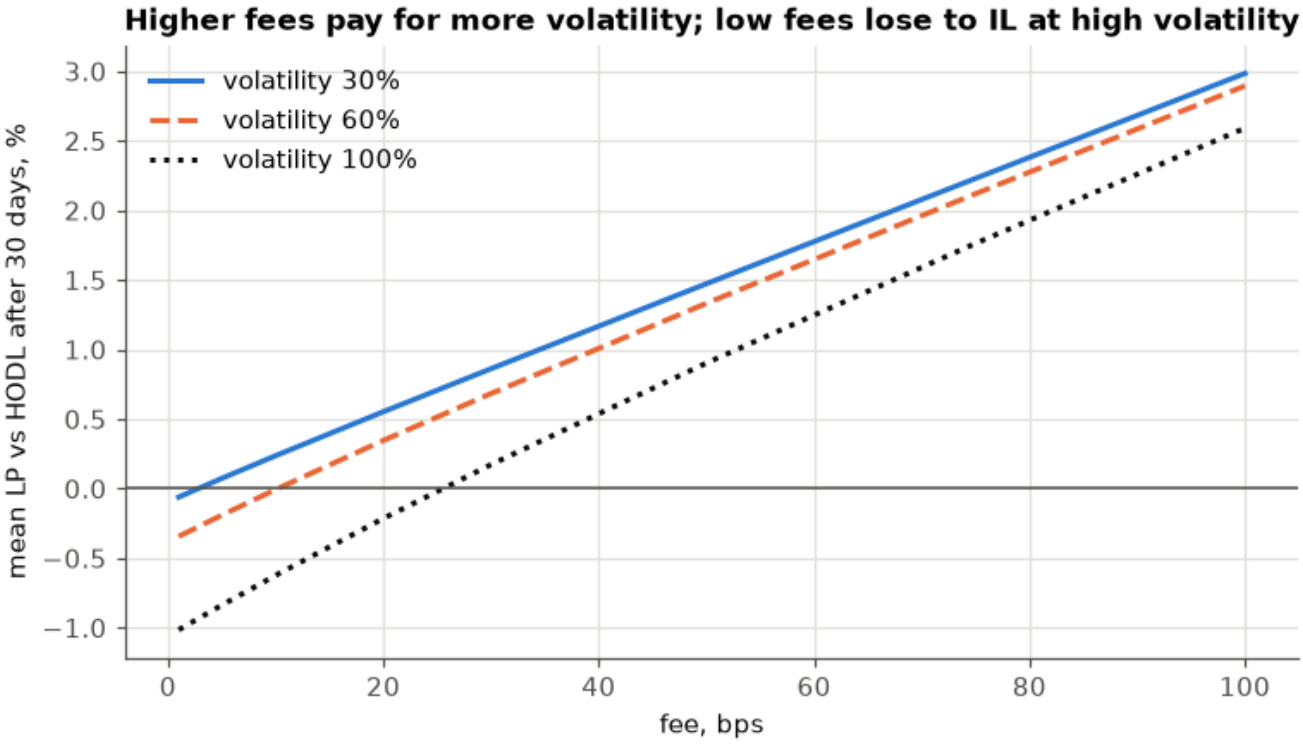


Figure 9. Higher fees pay for more volatility; low fees lose to IL at high volatility.

4.5 Reconciliation: model versus contracts

- **Local, done.** `scripts/exportSliceFixture.ts` runs the real `Pool`, `PriceGuard` and `OrderBook` on a local chain: a 20 WBTC order in 5 slices with and without refill, and the same 20 WBTC as one swap. The notebook rebuilds every slice amount, slice output, reserve after each slice, arbitrage trade and total from the starting reserves, and asserts equality: **56 comparisons, all exactly equal** (lump sum 1,150,517,367,388 raw USDC; 5 slices no refill 1,150,465,094,696; with refill 1,184,117,434,876).

- **Sepolia, done: 3 of 3 slices exactly equal.** Order 2 sold 5 WBTC in three slices (60 s apart, 3% tolerance). It was created from the web app and its slices were executed by our keeper script, with our arbitrage bot refilling the pool between them. For each `executeSlice` transaction, section 5 of `docs/analysis_slicing.ipynb` reads the `SliceExecuted` event and the pool reserves at the block before it, applies the integer swap formula, and asserts that the model's output equals the chain's. The reads are bundled in `docs/data/sepolia_order2_slices.json` so the notebook runs offline, and the cell reads the chain live when `SEPOLIA_RPC_URL` is set.

Slice	Transaction	Reserves before (raw WBTC / raw USDC)	Slice in (raw WBTC)	Model output = chain output (raw USDC)
1	`0x5e2d5767...`	49,906,524,485 / 38,923,228,585,325	166,666,666	129,167,076,412 (equal)
2	`0xb3ba7ea0...`	49,907,026,754 / 38,923,614,180,343	166,666,666	129,167,060,363 (equal)
3	`0x24a0d1fa...`	49,907,525,091 / 38,924,002,850,533	166,666,668	129,167,066,208 (equal)

What the order earned, compared with the same 5 WBTC sold at once from the pool state before slice 1 (the last two rows are model replays from that state, not other transactions):

Strategy	Proceeds (USDC)	vs lump sum
One lump-sum swap (replay)	384,946.34	n/a
Three slices, no refill (replay)	384,942.53	-0.1 bps
Three slices with the arbitrage refill (actual, on chain)	**387,501.20**	**+66.4 bps (+2,554.87 USDC)**

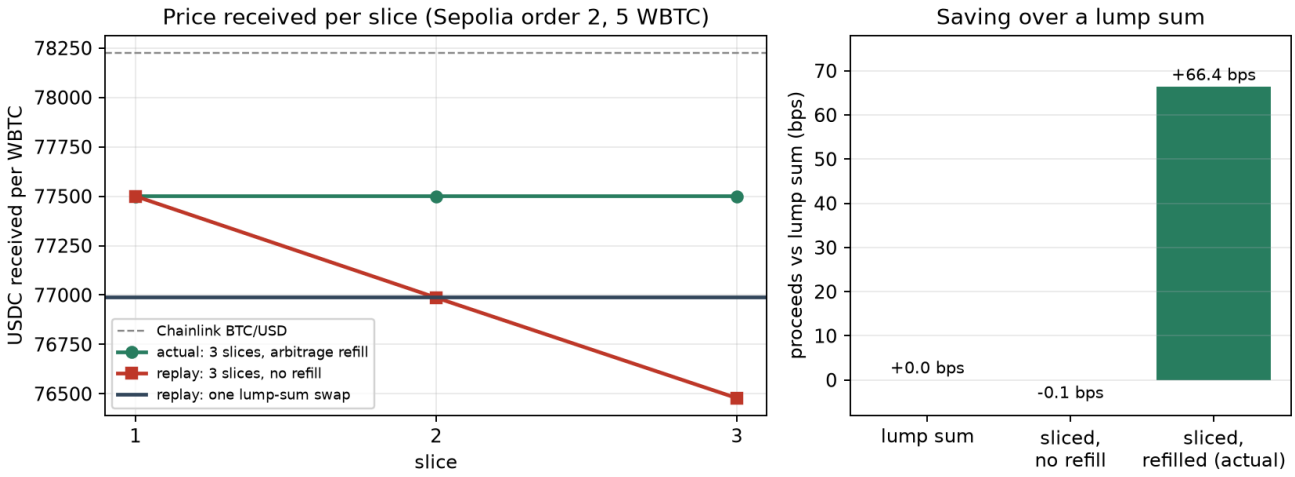


Figure 10. Sepolia order 2. Left: price received per WBTC for each slice. Right: proceeds against a lump sum.

This is the central claim of the project on a real network: slicing helped only because something refilled the pool (the no-refill replay is slightly *worse* than a lump sum), and the model predicted each slice to the last unit. One order is one data point: the Chainlink price rose between seeding the pool and the order, and the +66 bps depends on this pool's depth (about 500 WBTC for a 5 WBTC order).

4.6 What the model leaves out

Gas and keeper costs (measured separately in `docs/gas-orders.md`: about 160k to 200k gas per slice locally and about 510k on Sepolia; Section 5.3), the cost of the arbitrage that provides the refill and the competition for it, MEV on the slices (Section 2.4), a changing external price during the order (held constant here), correlation between volume and volatility, and a protocol fee. The break-even volume for protocol revenue promised in the proposal is **not estimated** because the fee is not implemented. All of these would make real outcomes worse than the numbers above.

4.7 Cost per user action and break-even (Part C)

Assumptions, stated once. Gas price 10 gwei on Ethereum L1, ETH at $3,000, BTC at $60,000. These are round assumptions, not observations; the notebook shows 1 to 50 gwei and $2,000 to $5,000. An "L2-like" column at 0.05 gwei shows how much of the cost is the chain rather than the design (rollup data fees are ignored). Gas was measured on a local Hardhat chain with the deployed compiler settings; **the Sepolia deployment contradicts the assumption that other networks cost the same.** Measured on Sepolia (2026-11-01, `docs/data/evidence_sepolia.json`): contract creation used about 7x the gas of our local figures (Pool 12.0M, OrderBook 15.7M, PriceGuard 3.9M, each mock token 4.2M; locally the three order-side contracts took 4.63M in total), a first-time `approve` estimates at 129,365 gas against 46,353 locally, and the seed `addLiquidity` used 925,874. The cause is not verified (new storage appears to be priced higher on that network). We therefore show both columns below and did not recompute every USD table. The USD figures below are therefore lower bounds for Sepolia-like pricing.

Measured on Sepolia (rehearsal order, 1 WBTC in 3 slices, 2026-11-01), against the local figures used in the table below:

Action	Local gas	Sepolia gas	Ratio
Approve WBTC (exact amount)	46,353	128,273	2.8x
`createOrder`	158,734	608,006	3.8x
`executeSlice` (first slice)	190,399	509,033	2.7x
`swapExactIn` (arbitrage refill)	117,775	234,251	2.0x
`claim`	56,705	60,164	1.1x

Actions that create new storage cost far more there; `claim`, which only updates existing storage, barely changes. One order each, so treat the ratios as indicative. The one-off deployment figure further down (4.63 M gas) is a local figure; on Sepolia the same three contracts used 31.7 M.

What each action costs.

Who pays	Action	Gas	USD at 10 gwei	USD at 50 gwei	USD, L2-like
Seller	Approve WBTC (exact amount)	46,353	1.39	6.95	0.007
Seller	`createOrder`	158,734	4.76	23.81	0.024
Seller	`claim`	56,705	1.70	8.51	0.009
Seller	`cancel` (unfilled)	56,148	1.68	8.42	0.008
Keeper	`executeSlice` (mean of 4)	163,545	4.91	24.53	0.025
Keeper	`executeSlice` (first slice, worst case)	190,399	5.71	28.56	0.029
Trader / arbitrageur	`swapExactIn`	117,775	3.53	17.67	0.018
LP	`addLiquidity` (existing LP)	123,548	3.71	18.53	0.019
LP	`removeLiquidity`	111,431	3.34	16.71	0.017

A full 10-slice order (approve, create, 10 slices, claim) is 1,897,242 gas: **$56.92** at 10 gwei and $3,000, from $3.79 (1 gwei, $2,000) to $474 (50 gwei, $5,000). The seller pays about $7.85 of that; the keeper pays the other $49.06. Deploying Pool, PriceGuard and OrderBook is 4.63 M gas, $139 once. Our numbers agree with the before/after tables in `gas-orders.md` and `gas-pool.md` to within a few percent; the differences come from account state, not the code.

Does slicing still pay after the keeper's gas? In the 500 WBTC / 30 M USDC pool, assuming arbitrage fully refills the pool between slices (the best case; Section 4.2 shows how partial refill shrinks it):

Order	Slices	Saving vs lump sum, before gas	Keeper gas (10 gwei)	Net saving	Net, bps of order
1 WBTC ($60k)	10	$107	$49	$58	9.7
1 WBTC ($60k)	20	$113	$98	$15	2.5
5 WBTC ($300k)	10	$2,655	$49	$2,606	86.9

Order	Slices	Saving vs lump sum, before gas	Keeper gas (10 gwei)	Net saving	Net, bps of order
20 WBTC ($1.2M)	10	$41,130	$49	$41,081	342.3
50 WBTC ($3M)	20	$256,331	$98	$256,233	854.1

Price impact grows with the square of the order's share of the pool, while the keeper's cost grows only with the number of slices. So gas only decides the question for small orders: on L1, slicing into 10 beats a lump sum from about **0.7 WBTC at 10 gwei** (1.6 WBTC at 50 gwei, 0.1 WBTC on an L2-like chain). Below that, more slices make it worse, not better. A cross-check against the real contracts: the end-to-end test's 20 WBTC, 5-slice order with our real arbitrage bot earned 292 bps over the lump sum, just under the model's full-refill ceiling of 303 bps for the same order.

Break-even for going to market. The proposal's revenue model has two parts, and **neither is built** (the contracts charge no execution fee and pay the whole 30 bps swap fee to LPs):

- a protocol share of the swap fee, assumed **5 of the 30 bps** (LPs keep 25), earned on the streaming volume and on the arbitrage volume that refills it (about equal to it: one refill per slice in our test);
- an **execution fee of 10 bps** of each order's notional, which also pays the keeper's gas.

Fixed costs per month, assumed: RPC $50, keeper and bot server $20, front-end hosting $20, monitoring $25, and an external audit of $40,000 amortised over 24 months ($1,667): **$1,782 in total**, no salaries. A second scenario adds one part-time engineer at $5,000 ($6,782).

Average order	Slices	Chain	Orders per month to break even	Streaming volume per month
$50,000	10	L1, 10 gwei	35	$1.75 M
$50,000	20	L1, 10 gwei	952	$47.6 M
$50,000	10	L1, 50 gwei	never (gas exceeds the fee)	never
$250,000	**10**	**L1, 10 gwei**	**4**	**$1.0 M**
$250,000	20	L1, 50 gwei	191	$47.8 M
$250,000	10	L2-like	4	$1.0 M
$1,000,000	10	L1, 10 gwei	1	$1.0 M

With one part-time engineer the base case ($250k, 10 slices, 10 gwei) needs **16 orders and $4.0 M a month**.

Three conclusions. First, **as built the protocol never breaks even**: it earns nothing, and someone (us, or a volunteer keeper) pays about $5 of gas per slice on L1. The execution fee is not optional for a real launch. Second, the business only works for **large orders**, which is also where slicing saves the most: a $250k order saves the seller tens to hundreds of bps and costs the seller only the 10 bps execution fee on top of the 30 bps swap fee it pays anyway. The protocol earns 20 bps of the order ($500): the execution fee plus its 5 bps cut of the swap fee on the order and on the arbitrage trades that refill it, both taken out of the LPs' 30 bps, not added to it. Small orders on a busy L1 lose money for the operator because 10 to 20 keeper transactions cost more than the fee; a minimum order size, or a per-slice fee floor, fixes that. Third, the levers are roughly equal: dropping the protocol share to 0 or the execution fee to 0 each takes the base case from 4 to 9 orders a month, so either one alone could carry a small operation.

What this leaves out. Gas and ETH prices and the order flow are assumptions, not data; break-even says how many orders are needed, not how many we would get. Full refill is the best case. Salaries, legal and marketing are excluded from the base case. Charging either fee needs contract changes (a fee switch in `Pool`, a fee in `OrderBook`) that were not made; **Decision:** the report presents both fees as future work (Sections 2.8, 6.4 and 7).

5. Testing, security and gas

5.1 Tests

One command, `yarn test` from the repository root, runs all **92 tests** (exit 0, run again on 2026-11-02 after a clean install).

File	Tests	What it covers
`test/Pool.ts`	24	liquidity round trips, hand-computed swap values, reverts, invariant fuzzing, TWAP
`test/MockTokens.ts`	9	decimals, owner mint, faucet cap and its revert
`test/PriceGuard.ts`	16	zero, negative, stale and future-dated prices, `answeredInRound`, the band at its exact limit, hard caps, non-owner reverts
`test/OrderBook.ts`	24	slicing and the last-slice remainder, claim after transferring the NFT, cancel, `TooEarly` revert, skip, reentrancy through an NFT receiver and through the pool
`test/EndToEnd.ts`	17	our real keeper and arbitrage code against the real contracts: streamed proceeds beat a lump sum (1,184,117 against 1,150,518 USDC, +292 bps), `TooEarly` revert, stale oracle, no double-send
`test/GasPool.ts`, `test/GasOrders.ts`	1 + 1	exact receipt gas for the before/after tables

Pool tests in detail.

- **Liquidity round trip.** Adding then removing all shares returns the deposit minus rounding and the 1000 locked shares (at least 99.9999% back). A second LP's round trip loses less than one share's worth of each token.
- **Hand-computed swap values.** Selling 1 WBTC into the 500 WBTC / 30,000,000 USDC pool returns exactly 59,700,956,293 raw USDC; selling 50 WBTC returns 2,719,832,681,640; selling 60,000 USDC returns 99,501,593 raw WBTC. These came from Python integer arithmetic, not from the contract.
- **Reverts.** Output below `minAmountOut` (including after a front-running trade moves the price), zero amounts on every function, an unknown token, an empty pool, a dust trade that rounds to zero, a first deposit too small for the locked minimum, and a TWAP window longer than the pool's history.
- **Invariant fuzzing.** A seeded loop of 250 random operations (swaps of up to 30% of a reserve in both directions, plus deposits and withdrawals) checks after every swap that `k = reserve0 * reserve1` has not decreased, and after every operation that the pool's actual token balances equal its recorded reserves. A second test replays five swap sizes against an independent integer model.
- **TWAP after known price changes.** A scenario with a price drop at +100 s and a rise at +300 s checks the average at four window lengths (500, 400, 250 and 100 s), including a window that starts between two records. Expected values were computed in Python.

Test quality check. A passing test only counts if it can fail, so we changed the pool three ways and confirmed the suite caught each: fee changed from 30 to 25 bps (9 tests fail), TWAP interpolation removed (2 fail), slippage check removed (2 fail). The contract was restored and re-tested after each.

Limits of the testing. The fuzz test is a seeded TypeScript loop, not coverage-guided fuzzing (Foundry, Echidna), and 250 operations is a sample, not a proof. The pool was tested only with 8-decimal and 6-decimal tokens; it reads decimals generically but other pairs were not tried. Live Sepolia behaviour is covered by the transactions in Sections 3.5 and 4.5, not by these tests.

Order side: test quality check. We changed `OrderBook` four ways (last slice leaves dust, `claim` pays the wrong address, `nonReentrant` removed, `TooEarly` check removed) and confirmed the suite fails each time. Reverts are tested for slippage, zero amounts, unknown tokens, early slices, non-owners, an unusable oracle and finished orders.

5.2 Slither

Slither 0.11.4 with 100 detectors, on every contract file and on one combined compilation unit of `Pool`, `OrderBook` and `PriceGuard`: **no High-severity finding anywhere**. The combined unit has 20 results, each fixed or explained (`docs/slither-full.md`, with the per-contract detail in `docs/slither-pool.md` and `docs/slither-orders.md`):

Count	Detector	Impact	Verdict
1	reentrancy-no-eth (`OrderBook.executeSlice` writes `proceeds` after the swap)	Medium	explained: the amount exists only after the swap; every state-changing function is `nonReentrant` and the pool is immutable
3	divide-before-multiply (`Pool`)	Medium	explained: floor in the pool's favour; exact in the TWAP

Count	Detector	Impact	Verdict
6	incorrect-equality (`Pool`)	Medium	explained: array lengths and the empty-pool state; the pool never reads `balanceOf`
1	unused-return (`PriceGuard` ignores `startedAt`)	Medium	explained: `answer`, `updatedAt` and `answeredInRound` carry the safety information
7	timestamp	Low	accepted: scheduling and staleness are the features
2	solc-version and naming on the frozen interface files	Info	accepted

The pool findings in more detail: Slither 0.11.4 with 100 detectors, on the final code: 0 findings on both mock tokens and 16 results from 6 detectors on `Pool`, none of them a vulnerability and none about reentrancy. Every finding is fixed or explained in `docs/slither-pool.md`. In short: three divide-before-multiply notes (one exact by construction because the TWAP step is `price * seconds`, one rounding in the pool's favour, one a false positive from a branch Slither cannot tell apart); `== 0` checks on array lengths and reserves (safe, because the pool never reads `balanceOf`, which is the pattern that detector exists to catch); `block.timestamp` use (needed for a TWAP); and two notes about the frozen interface file. One finding led to a change: after we packed the reserves, the unused-return warning on the range check disappeared. The report also lists what Slither cannot see (no admin, first-depositor price setting, sandwiching, TWAP manipulation cost).

One change followed from the first Slither run: the implementation contracts used the scaffold pragma `>=0.8.0 <0.9.0` and now use `^0.8.20`. The scan covered the Scaffold-ETH sample contract, which we afterwards deleted (it was not part of the product); the production contracts did not change and Slither was not re-run after the deletion. Slither does not analyse the TypeScript scripts or the front end.

5.3 Gas

Gas was measured as exact receipt gas (solc 0.8.30, optimizer 200 runs). We made **two changes because of it**, and one attempt made things worse.

Pool: pack reserves and update time into one storage slot (112 + 112 + 32 bits, the layout `getReserves()` returns):

Function	Before (gas)	After (gas)	Saved
addLiquidity (first deposit)	288,208	244,666	15.1%
addLiquidity (existing LP)	132,852	123,548	7.0%
swapExactIn (sell WBTC)	114,761	105,419	8.1%
swapExactIn (sell USDC)	114,832	105,487	8.1%
removeLiquidity (half of an LP's shares)	120,694	111,431	7.7%

Order book: store the order in two slots instead of three (the interface struct stays unchanged; the stored slice size is narrowed to 96 bits):

Function	Before (gas)	After (gas)	Saved
createOrder (first order)	185,588	163,534	22,054 (11.9%)
executeSlice (first slice)	201,777	199,791	1,986 (1.0%)
cancel	60,153	56,148	4,005 (6.7%)

Our first attempt, caching the order in memory inside `executeSlice`, was **924 gas worse** and was reverted.

The same actions on Sepolia cost more. Measured from the receipts of the live transactions (rehearsal and demo orders, 2026-11-01 and 2026-11-02; `docs/data/evidence_sepolia.json`):

Action	Local gas	Sepolia gas	Ratio
Approve WBTC (exact amount)	46,353	128,273	2.8x
`createOrder`	158,734	608,006	3.8x

Action	Local gas	Sepolia gas	Ratio
`executeSlice` (first slice, by the keeper)	190,399	509,033	2.7x
`swapExactIn` (arbitrage refill)	117,775	234,251	2.0x
`claim`	56,705	60,164	1.1x
`addLiquidity` (first deposit, seeding the pool)	244,666	925,874	3.8x
Deploy `Pool` / `OrderBook` / `PriceGuard`	n/a	12,035,012 / 15,699,128 / 3,924,284	n/a

Actions that create new storage cost 2 to 4 times more than in our local tests, while `claim`, which only updates existing storage, is about the same. We have not established why the network prices storage creation higher, and it was a surprise to us: our first cost estimates assumed Sepolia would match the local EVM. The deployment cost about 0.036 ETH, mostly because the deploy tool used a fixed 1 gwei priority fee while the network was suggesting about 0.0005 gwei; the bots, which ask the RPC for its suggested fee, pay almost nothing.

5.4 What the testing does not cover

- AI assistance was used for parts of the code, tests and documentation (Section 8); all of it was checked by the test suite, mutation checks, Slither and independent Python models.
- No external audit, no fuzzing beyond the seeded loop in `test/Pool.ts`, and no test of an ERC-20 with other decimals than 8 and 6.
- The front end was driven in a browser against a local chain and against Sepolia with a real wallet for approve and create; claim was sent from the same wallet, and cancel and the stale-feed warning were not exercised on Sepolia.

6. Going to market

6.1 Users

The people who lose the most to price impact are those who must sell a lot at once and cannot wait for a perfect market: **DAO and company treasuries** paying expenses in stablecoins, **funds** rebalancing or meeting redemptions, **miners and long-term holders** converting BTC exposure, and **token teams** selling grants or unlocks. Today they choose between eating slippage on a public pool, splitting the trade by hand (which needs someone watching for hours), or an OTC desk (a quoted spread, counterparty trust, and minimum sizes). Our order runs unattended, on public rails, and the seller can see every slice.

Part C (Section 4.7) shows who we are actually for: on Ethereum L1 slicing pays after gas from about **0.7 WBTC** (~$40k) at 10 gwei, and the saving becomes large only when the order is a meaningful share of the pool (342 bps for a 20 WBTC order in a 500 WBTC pool with 10 slices and full refill). Retail-size trades should keep using a normal swap.

6.2 Trust

A treasury will not hand a large order to an unknown contract on our word. What we can show instead:

- **Non-custodial ownership.** The order is an NFT held by the seller. Only the holder can cancel, and cancelling returns all unsold WBTC and unclaimed USDC in the same transaction (tested, including after transferring the NFT).
- **Almost no admin.** `Pool` and `OrderBook` have **no owner at all**: no pause, no upgrade, no fee switch, no way to move user funds. `PriceGuard` has one owner who can change exactly two numbers, each with a hard cap and an event: the staleness limit (at most 1 day) and the allowed deviation (at most 20%). It cannot change the feed or touch funds. This list is in the source comments and in Section 2.6.
- **No trading blind.** If Chainlink is stale, zero or negative, or the round is carried over, slices revert; if the pool is far from Chainlink, they are skipped. Every slice must return at least the Chainlink-implied amount minus the seller's own tolerance. The web app shows a warning as soon as the feed is stale or close to it.

- **Evidence.** Verified source on Etherscan (Section 3.5), Slither results with every finding fixed or explained (`docs/slither-full.md`: no high-severity finding), 92 automated tests including expected-revert, reentrancy, a seeded invariant test on the pool, and mutation checks; gas reports; and notebooks that reproduce the contracts to the last unit.
- **What we cannot offer yet.** No external audit, no track record, a test deployment only. A real launch needs an audit first (budgeted in Part C) and a capped period with small limits.

6.3 The first hundred users

1. **Ten design partners first.** DAO treasuries sell in public: their governance forums list planned sales. We would offer the first ten a free, hand-held order (we run the keeper and the arbitrage for them) in exchange for a public write-up. Ten real orders are enough to publish a first comparison.
2. **A public fill dashboard.** For every order, show what it received against a lump-sum swap at the moment it was created, computed with the same model the notebooks use. Proof on chain is our main marketing; the web app's preview table is the start of it.
3. **Liquidity first.** Slicing only works if the pool is deep and refilled, so the first money goes to LPs and arbitrageurs, not sellers: early-LP rewards (the proposal suggested paying them in our share token; that requires a reward contract we have not built; future work) and an open-source arbitrage bot, which we already have.
4. **Distribution through existing tools.** A Safe (multisig) app, since treasuries sign through Safe, and integration with aggregators so a large swap can be routed into a streaming order automatically.
5. **Keepers as a community.** `executeSlice` is permissionless. With the execution fee (6.4) part of it can be paid to whoever runs the slice, so other people run keepers and the system does not depend on us.

6.4 Business model with numbers

Revenue would come from a **protocol share of the swap fee** (assumed 5 of the 30 bps; LPs keep 25) and an **execution fee** (assumed 10 bps of each order's notional, which also pays the keeper's gas). **Neither is implemented in the contracts: both are future work**, and as built the protocol earns nothing. Each needs a contract change (Section 4.7). With fixed costs of about $1,800 a month (infrastructure and an audit amortised over two years, no salaries), at 10 gwei and $3,000 ETH, Part C finds:

Average order	Break-even orders / month	Monthly streaming volume
$50,000, 10 slices	35	$1.75 M
$250,000, 10 slices	4	$1.0 M
$1,000,000, 10 slices	1	$1.0 M

With one part-time engineer ($6,800 a month) the $250k case needs 16 orders and $4 M a month. For comparison, a seller of a $250k order in our test pool saves far more than the 10 bps execution fee (Section 4.7). Two rules follow: a **minimum order size** (or a per-slice fee floor), because on a busy L1 a small order's 10 to 20 keeper transactions cost more than its fee; and a **cheap chain** for smaller orders, where keeper gas is a few cents.

6.5 Risks

Risk	What could happen	What we do or would do
Smart-contract bug	Funds locked or lost. The slice logic is the likeliest place.	Fixed slice count, flat fee, no admin, reentrancy guards, 92 tests, Slither; an external audit before any real money.
Oracle failure	Chainlink stops updating or reports a wrong price.	Stale or invalid prices stop trading; deviation band; per-order tolerance. A wrong-but-fresh price is not detected.
No refill	Nobody arbitrages, so slicing saves nothing and the guard skips slices.	Our own open-source bot; liquidity incentives; the app shows the "no refill" estimate next to the best case.
MEV / sandwich	A searcher moves the pool before a slice.	Loss per slice is bounded by the deviation band and the tolerance. Not prevented (no private mempool). (Section 2.4)

Risk	What could happen	What we do or would do
Keeper downtime	Slices run late.	Permissionless execution; the seller can run a slice from the app; the next slice is always one interval after the last attempt, so a late keeper never causes a burst.
Thin liquidity	Large orders still move the price a lot.	Preview before signing; minimum interval; deeper pools come first.
Competition	CoW Swap's TWAP orders, Uniswap X, TWAMM-style hooks and OTC desks already serve large sellers.	Our edge is the simplest possible trust model (no owner, NFT ownership, oracle guard) and published fill evidence. Honestly, it is a small edge.
Regulation	Running a front end and collecting fees may bring obligations.	Legal review before charging any fee. Not assessed here.

7. Summary

We built and tested a streaming-order AMM: a 30 bps constant-product pool with LP shares and a TWAP, a Chainlink price guard, and an ERC-721 order book that sells a large order in equal slices, plus a keeper, an arbitrage bot and a web app.

What we found. Slicing saves money only when the pool is refilled between slices. Without refill, 20 WBTC sold in 5 slices returns 1,150,465 USDC against 1,150,517 for one swap, so slightly less, because each slice pays the fee on a worse price. With our arbitrage bot refilling the pool, the same order returned 1,184,117 USDC, **292 bps more**, on the real contracts. The notebook's integer model reproduces the contracts exactly (56 of 56 comparisons locally). **On Sepolia** the same formula, applied to the pool reserves read from the block before each slice, gives the on-chain output to the last unit for all three slices of a 5 WBTC order (slice 1 `0x5e2d5767...`). That order, created from the web app and filled by our keeper with our arbitrage bot refilling the pool, returned 387,501.20 USDC, against 384,946.34 for the same 5 WBTC as one swap from the same pool state: **+2,554.87 USDC, +66 bps**; replayed without the refill, the three slices would have returned 384,942.53, 3.80 USDC less than the lump sum. After gas, slicing pays on L1 from about 0.7 WBTC at 10 gwei, and the saving grows quickly with order size. Without refill the price guard also stops a large order from finishing: a 50 WBTC order in 20 slices is first skipped at slice 7.

What it would take as a business. As built the protocol earns nothing. With a 5 bps share of swap fees and a 10 bps execution fee it covers about $1,800 a month of running costs with four $250k orders a month on L1. That is plausible, but only for large orders, and it requires contract changes we did not make.

What is solid and what is not. Solid: five contracts deployed and verified on Sepolia, a full order created, filled and claimed there, 92 passing tests, Slither with no high-severity finding, two measured gas improvements (pool swaps -8.1%, `createOrder` -22,054 gas), and exact model-to-contract agreement. Not yet: an external audit; protection against sandwiching; the proposed fees and LP incentives.

Steps we took.

1. Proposal (Oct 14): topic, contracts, features, split of work.
2. Scaffold-ETH 2 project; interfaces agreed and frozen; one price unit chosen (Section 2.1).
3. Pool and test tokens with tests, then PriceGuard and OrderBook with tests; Slither and the two gas changes (Section 5).
4. Analysis notebooks A and B, built on the same integer arithmetic as the contracts, and checked against the real contracts locally (56 of 56).
5. Keeper, arbitrage bot and end-to-end tests; then the web app and the cost and break-even analysis (Part C).
6. Sepolia: deployment and Etherscan verification (Nov 1), a rehearsal order, then an order created in the web app and filled by the keeper with the bot refilling the pool.
7. Reconciliation of the model with the three on-chain slices; this report and the video.

Unrealized profit and loss, under stated assumptions. Three positions in the system carry unrealized P&L before they are closed.

- **The seller's order.** Unsold WBTC is marked at the Chainlink price. The order we ran realized +2,554.87 USDC (+66 bps) over the same coins sold at once from the same pool state (Section 4.5; the lump-sum figure is a

model replay).

- **A liquidity provider.** Against simply holding the two tokens, the LP's unrealized P&L is fees minus impermanent loss. In the base case of Section 4.4 (30 days, simulated price paths, stated volume and volatility) fees average 1.07% of TVL against a mean impermanent loss of 0.38%, so the LP is on average 0.68% ahead of holding and ahead on 91.6% of paths. The assumptions (price process, volume, arbitrage timing) are in `docs/analysis_pool.ipynb`.
- **The arbitrage bot.** Each refill leaves the bot holding more WBTC. At the pool's own quote, valued at the Chainlink price with USDC at one dollar, its four refills around the demo order had expected profits of 445.09, 431.34, 431.36 and 431.35 USDC (1,739.14 in total). These are unrealized: they assume the inventory is valued at the oracle price, and they ignore gas, which was negligible on Sepolia.

Advanced functionality delivered. Oracle: a Chainlink guard (staleness, sign and round checks, plus a deviation band) and a TWAP accumulator in the pool. Front end: a full Scaffold-ETH 2 application with a confirmation step, order progress, claim and cancel, transaction status, plain-language errors and a stale-feed warning. Gas measured and optimized: two changes with before and after numbers, and the measured Sepolia cost of each action (Section 5.3). Also: an off-chain keeper and arbitrage bot, a one-command test suite with revert tests, and a Slither scan with every finding explained. Not done: a layer-2 deployment and a token promotion campaign.

What we learned.

- **The saving comes from the refill, not from slicing.** Splitting an order only reorders the sale; the price ends where the reserves put it. Everything the product delivers depends on someone moving the price back between slices, so the design question is who does that and what it costs.
- **Decimals are the bug that matters.** Pricing an 8-decimal and a 6-decimal token against an 8-decimal feed made the first design 100 times off; one price unit and one conversion library fixed it (Section 2.1).
- **A local test is not the network.** Sepolia surprised us on gas (2 to 4 times), on log queries (10 blocks on a free tier), on the feed's hourly update, and on the arbitrage threshold. Each of those became a change in Section 2.9.
- **A model that matches the chain is worth more than a good-looking number.** Replaying every slice against reserves read from the chain, and breaking the code on purpose to see the tests fail, found more than reading the code did.

8. AI Use

AI was used in the making of this project. We worked with an AI coding assistant (Claude Code, by Anthropic) to help write and revise the smart contracts, tests, off-chain scripts, web application, analysis notebooks and documentation, including this report.

We checked its output the same way we checked our own work: the full test suite (92 tests), deliberate changes to the code that the tests had to catch, a Slither security scan, Python models that reproduce the contracts' arithmetic exactly, and live transactions on Sepolia that the model matches to the last unit (Section 4.5). A log of how the assistant was used is kept in the repository (`docs/ai-use.md`). The team is responsible for the submitted work.

Appendix A. Execution evidence: Etherscan links and screenshots

Every contract and transaction below is on Ethereum Sepolia. Contract source is verified on Etherscan: open an address in Section 3.5 and choose the Contract tab.

A.1 Transactions

Transaction	What it shows
`0xa9f51895...`	Pool deployed (verified on Etherscan)
`0xe52450a3...`	PriceGuard deployed (verified)
`0xbb0acbd1...`	OrderBook deployed (verified)

Transaction	What it shows
`0x688d1fd6...`	MockWBTC deployed (verified)
`0xae7c03dc...`	MockUSDC deployed (verified)
`0xf2bea583...`	first addLiquidity: pool seeded at the live Chainlink price, LP shares minted
`0x9d02d4bd...`	createOrder from the web app: 5 WBTC, 3 slices, 60 s, 3% tolerance; order NFT minted
`0x5e2d5767...`	executeSlice 1 by the keeper script (slice output reconciled with the model to the last unit)
`0xb3ba7ea0...`	executeSlice 2 by the keeper script
`0x24a0d1fa...`	executeSlice 3 by the keeper script (last slice, remainder)
`0xaadb43e7...`	arbitrage bot trades the pool back toward Chainlink before slice 1 (refill)
`0x62d1c0a5...`	arbitrage bot refill between slices 1 and 2
`0x3253d705...`	claim by the NFT holder's wallet: 387,501.20 USDC paid to the holder
`0x6e434545...`	rehearsal order 1 created (1 WBTC, 3 slices)
`0x3a523cc9...`	rehearsal order 1, slice 1 by the keeper
`0x5e24556a...`	rehearsal order 1 claimed
`0xd41947f7...`	video demo order 5: createOrder from the web app (2 WBTC, 3 slices, 60 s)
`0xa995edcc...`	video demo order 5: slice 1 by the keeper script
`0xc03addd9...`	video demo order 5: claim of the first slice's proceeds (50,982.31 USDC)
`0x0b4a33c0...`	video demo order 5: slice 2 by the keeper script
`0xdd856a6b...`	video demo order 5: slice 3 (last) by the keeper script

A.2 Keeper log (order 2, the order created in the web app)

Three consecutive lines from `yarn keeper`, copied unchanged. Each carries the transaction hash of a slice the keeper sent.


```
{"ts":"2026-11-
01T23:55:02.069Z","source":"keeper","event":"filled","orderId":"2","txHash":"0x5e2d57674c0b396c4c36d080436d7d655d7
cc986467b3a595591e76c3f63afa7","detail":"slice executed, 2
left","amountIn":"166666666","amountOut":"129167076412","slicesLeft":2}
{"ts":"2026-11-
01T23:56:13.598Z","source":"keeper","event":"filled","orderId":"2","txHash":"0xb3ba7ea07ecbee1d9ed18f651070cb56814
3cf43b1d73827bb78ef25c27ac942","detail":"slice executed, 1
left","amountIn":"166666666","amountOut":"129167060363","slicesLeft":1}
{"ts":"2026-11-
01T23:57:28.756Z","source":"keeper","event":"filled","orderId":"2","txHash":"0x24a0d1fae617dd5ddd4482a2e74e3875edd
d3df0a4de36f908a9c35700fe297c","detail":"slice executed, 0
left","amountIn":"166666668","amountOut":"129167066208","slicesLeft":0}
```

A.3 Screenshots

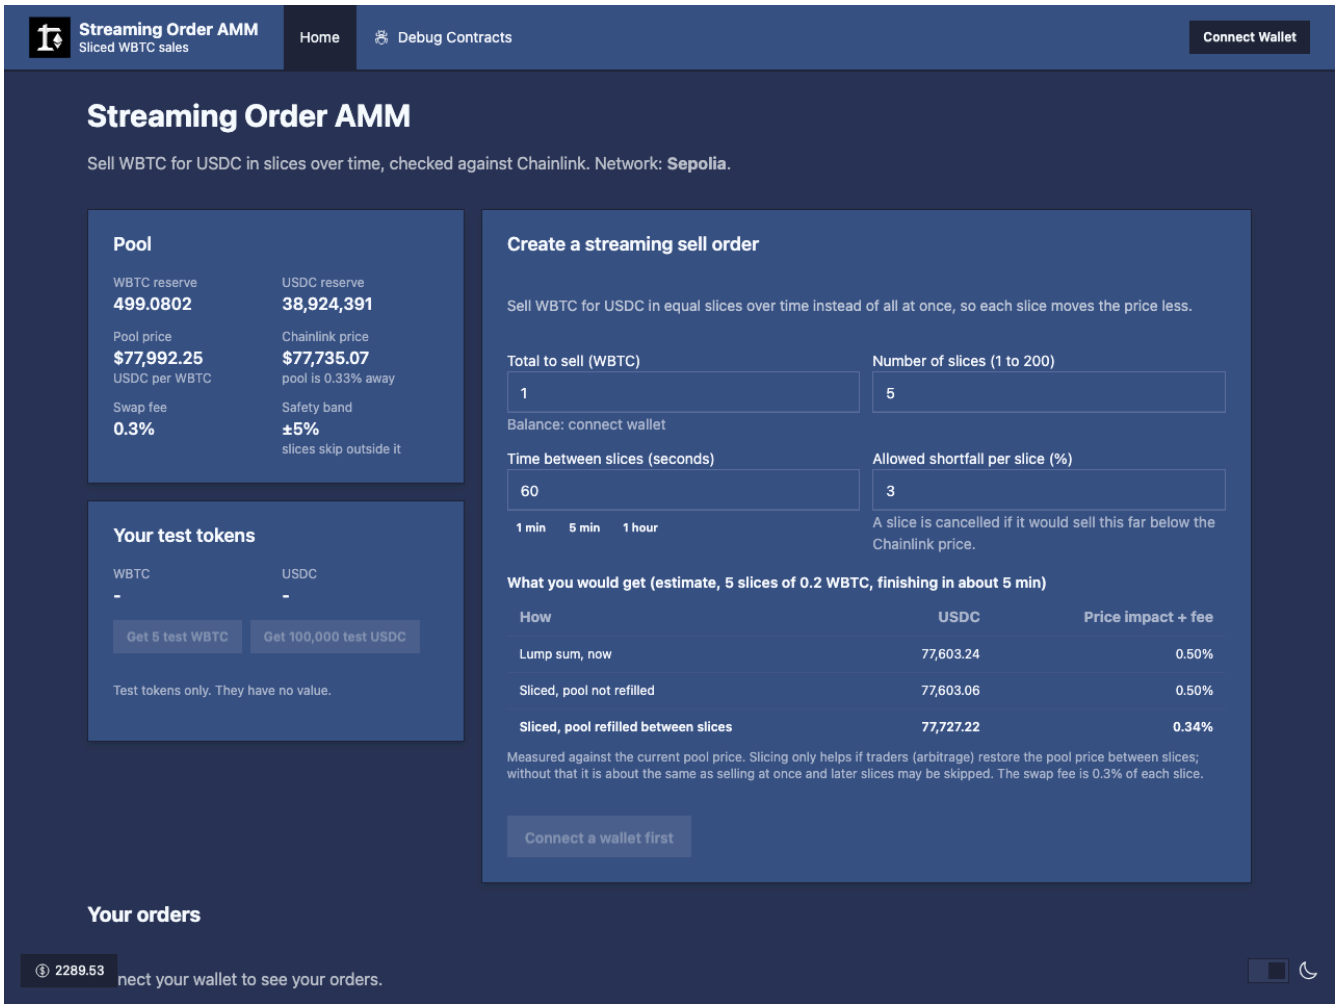


Figure 11. Web application on Sepolia: live pool state, Chainlink price, the create-order form and the lump-sum versus sliced preview.

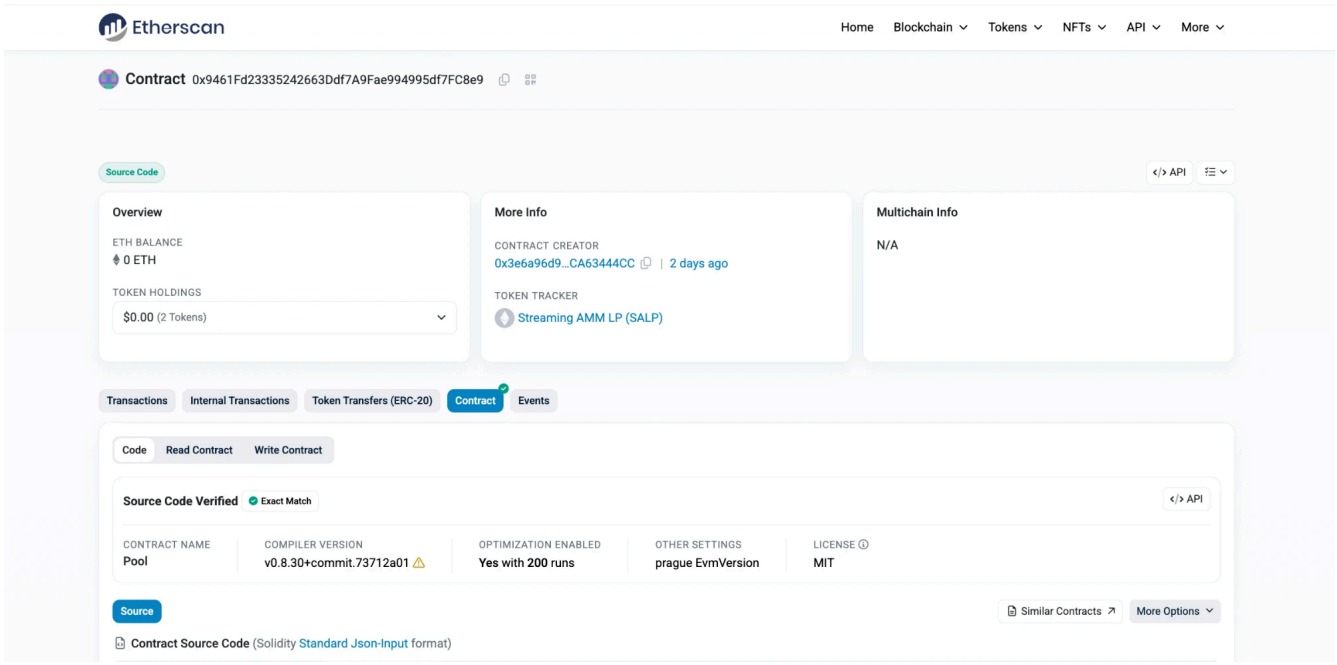


Figure 12. Pool on Sepolia Etherscan, source verified.

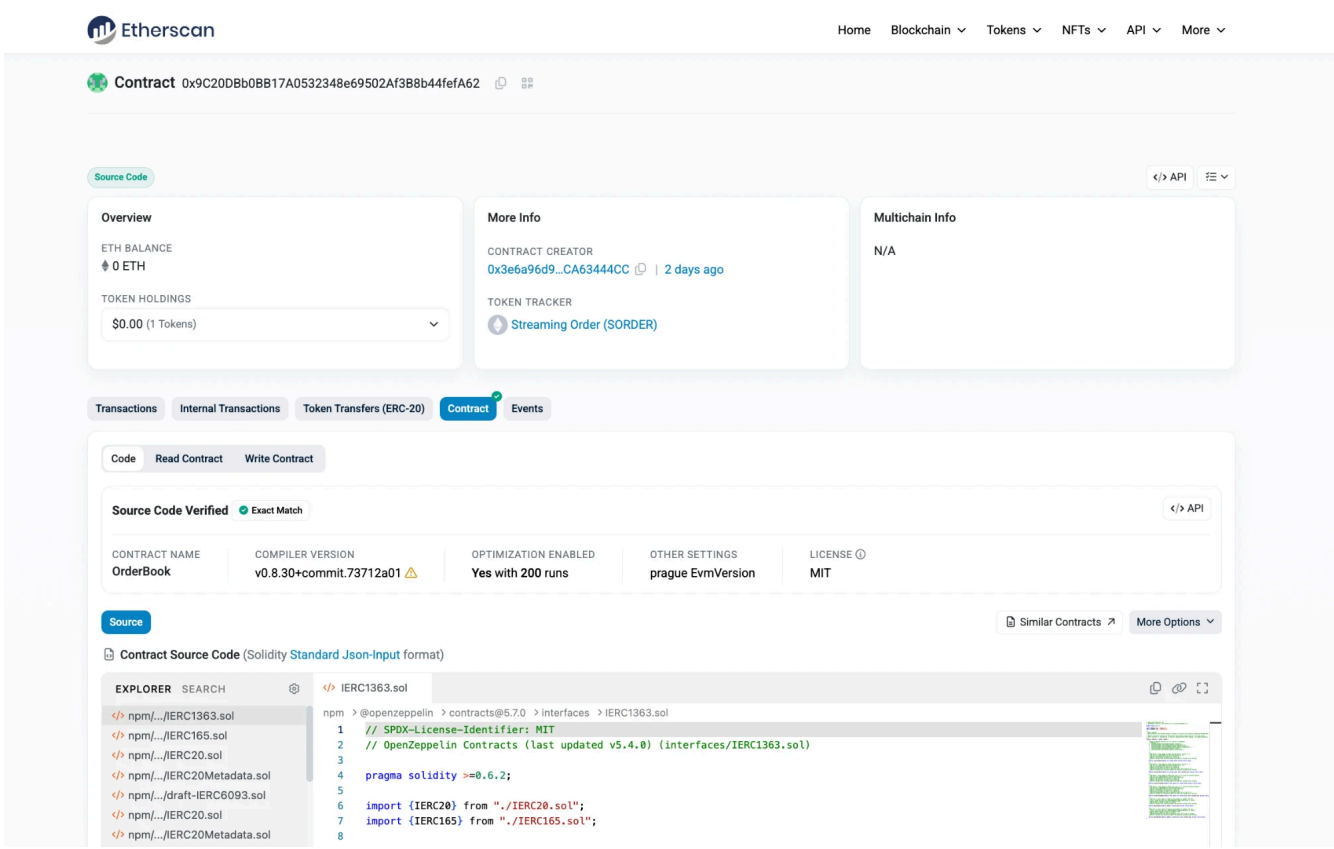


Figure 13. OrderBook on Sepolia Etherscan, source verified.

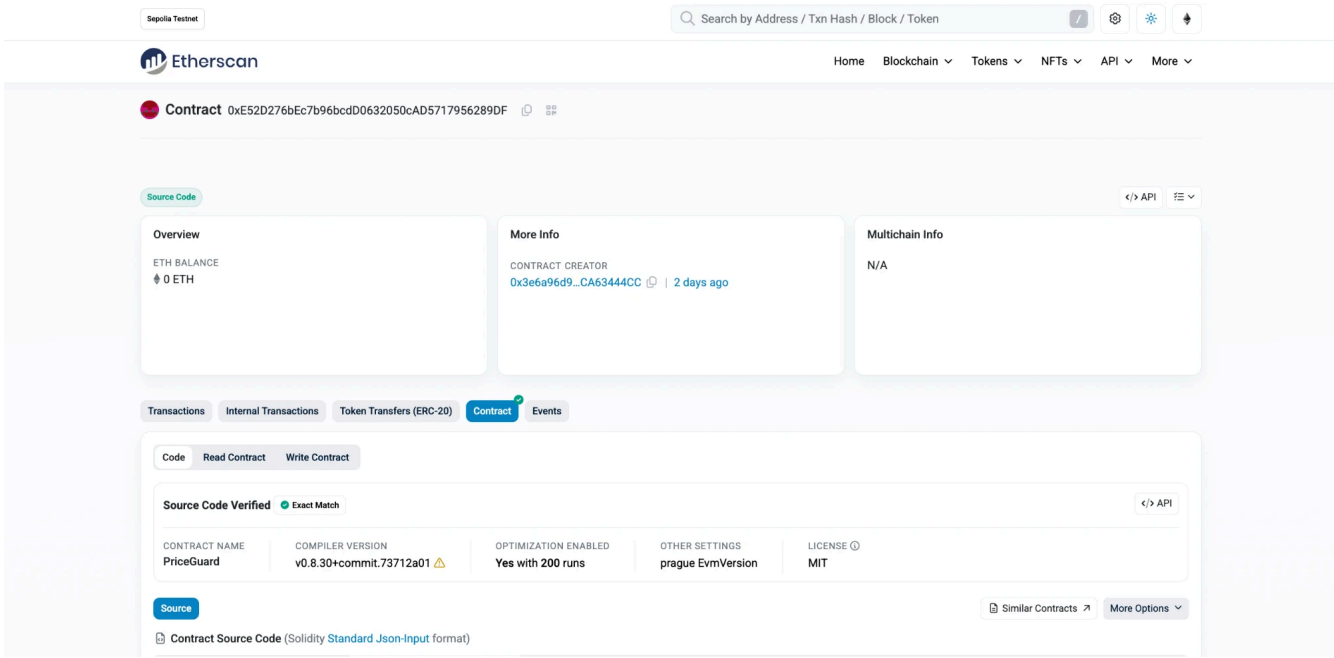


Figure 14. PriceGuard on Sepolia Etherscan, source verified.

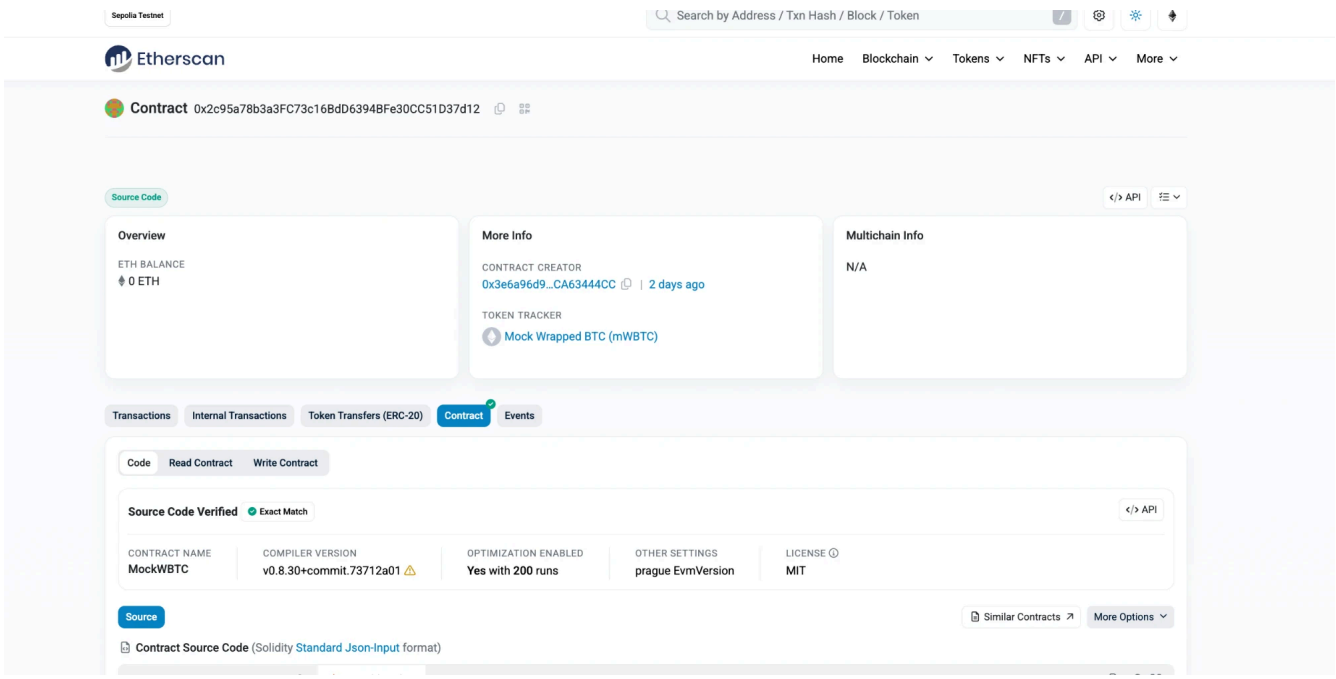


Figure 15. MockWBTC on Sepolia Etherscan, source verified.

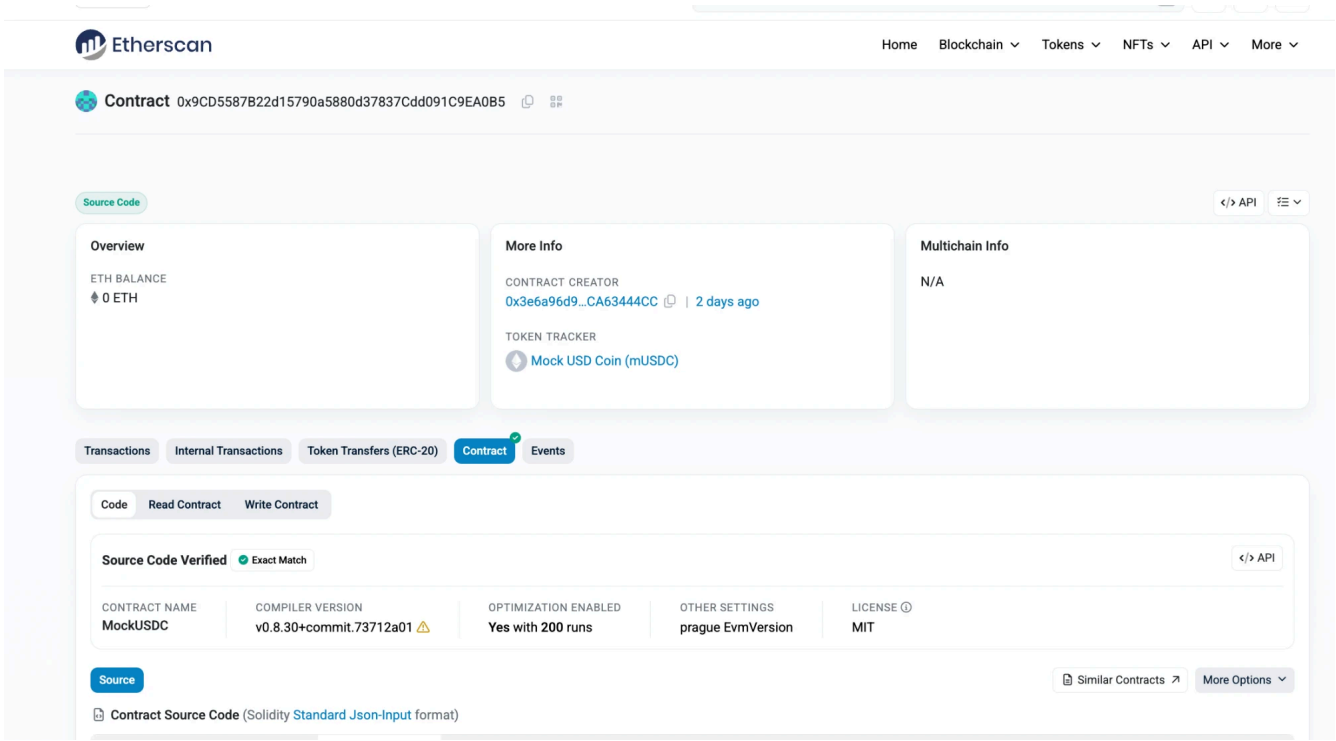


Figure 16. MockUSDC on Sepolia Etherscan, source verified.

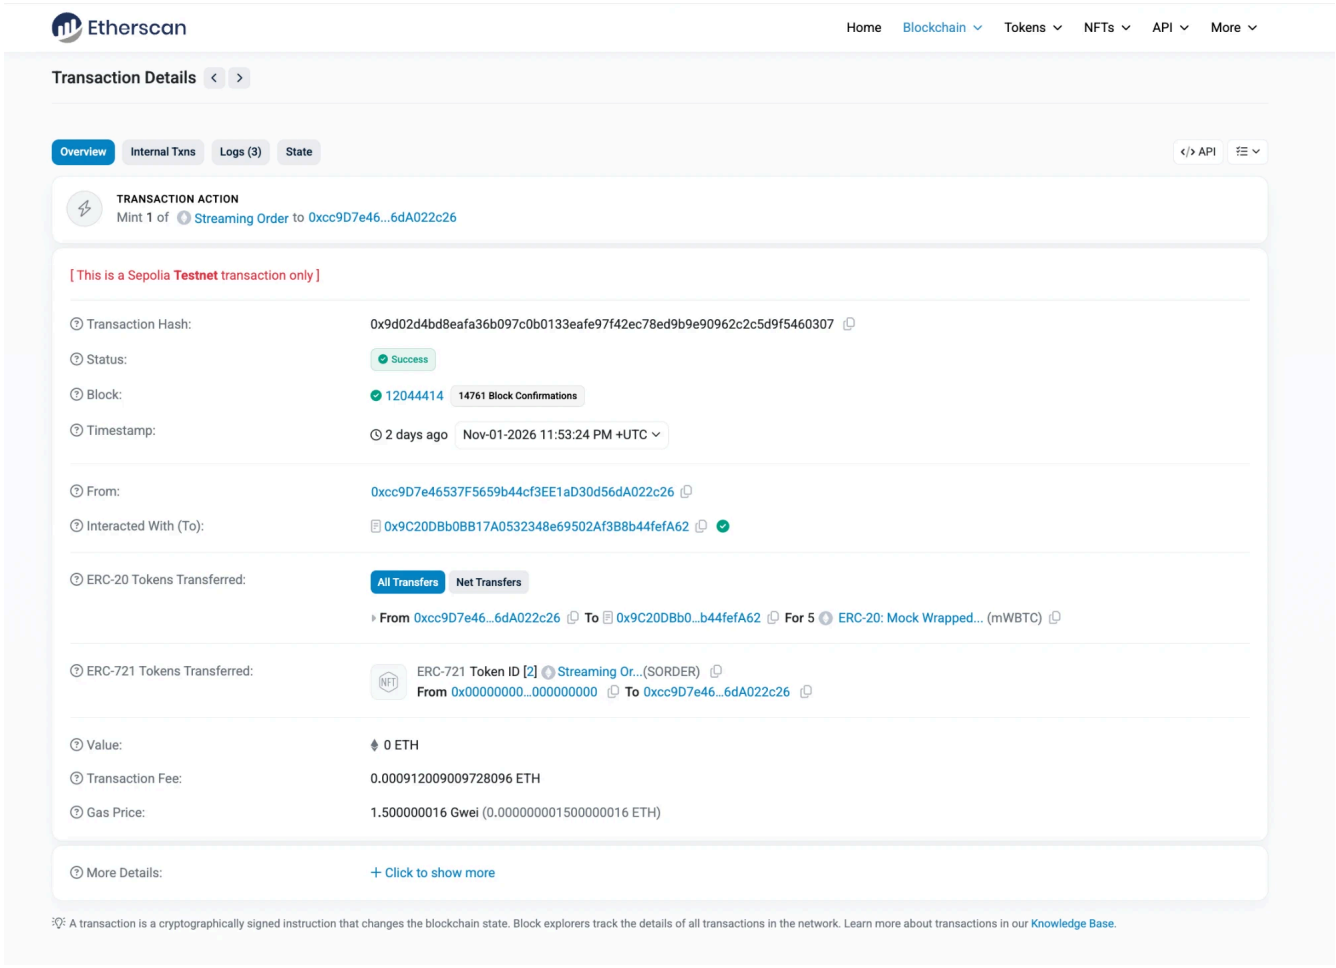


Figure 17. createOrder transaction (order 2) on Etherscan.

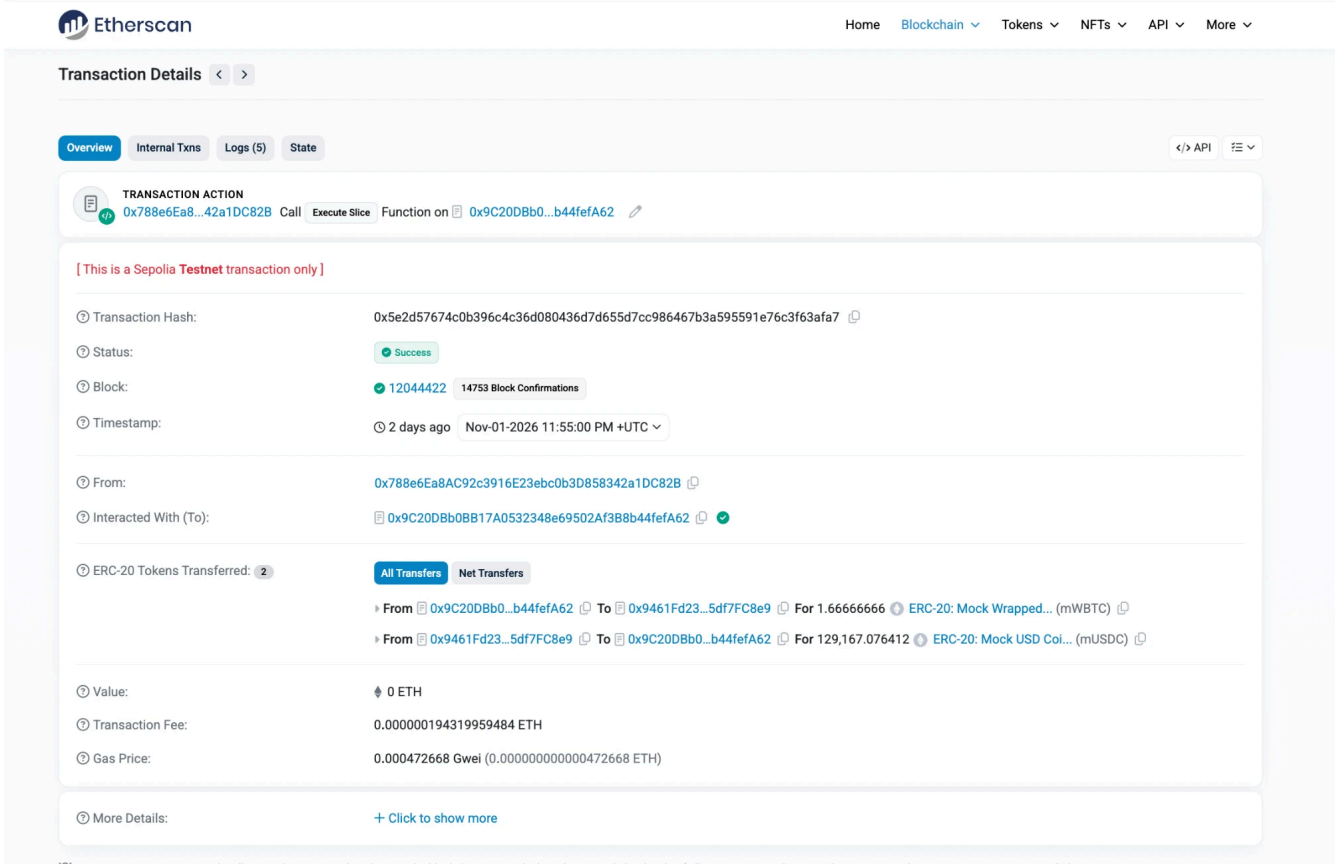


Figure 18. executeSlice transaction by the keeper on Etherscan.

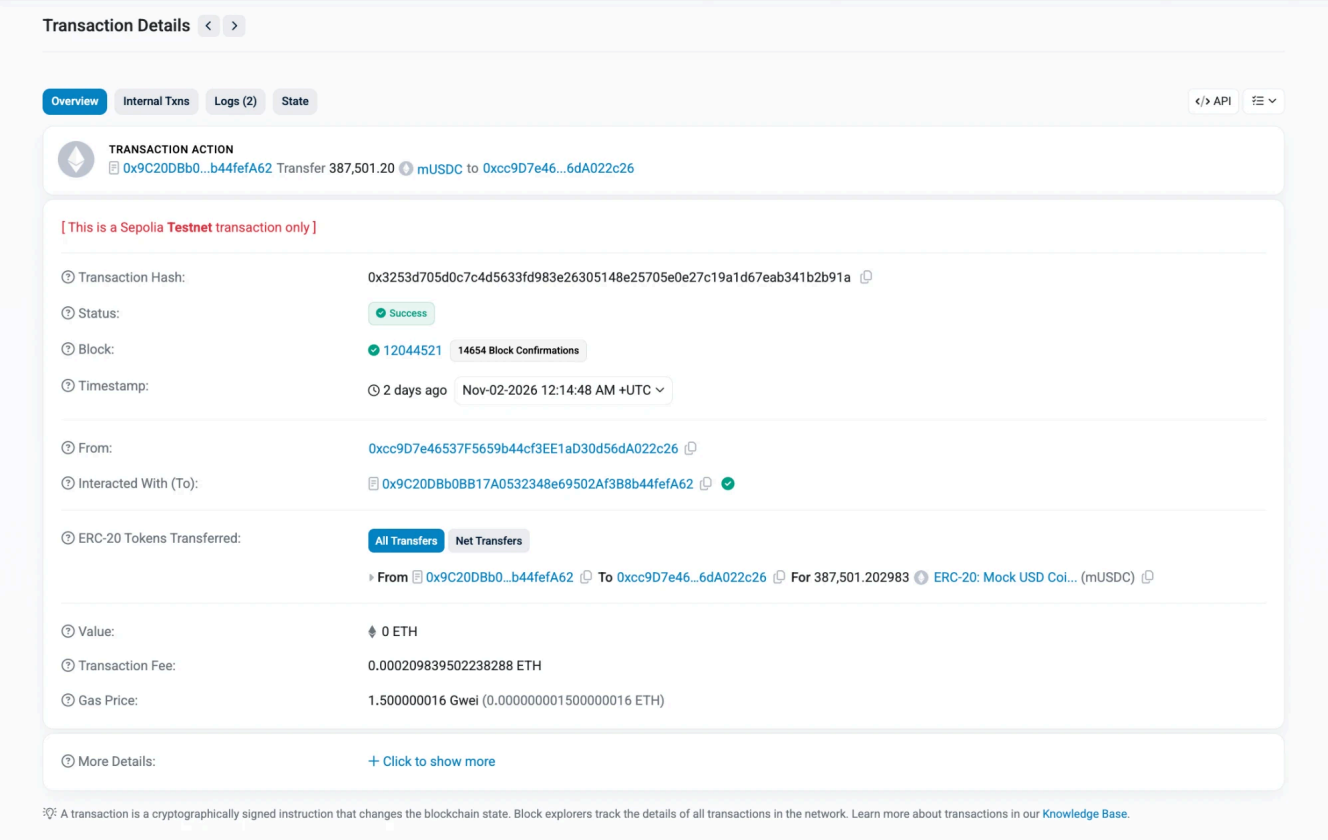


Figure 19. claim transaction on Etherscan.

Appendix B. Source code, metadata and JSON

- **Solidity source:** `docs/source-code.pdf` lists every production contract with line numbers (Pool, OrderBook, PriceGuard, PriceMath, the three interfaces and the two test tokens), and the same files are in `packages/hardhat/contracts/`. Public and external functions carry NatSpec.
- **Off-chain code:** `packages/hardhat/scripts/` (keeper, arbitrage bot, evidence and funding tools) and `packages/nextjs/` (web application).
- **Metadata and JSON:** `submission.json` (team, token, contracts, transactions, features); `packages/hardhat/deployments.json` (addresses, deployment transactions, blocks, constructor arguments); `packages/hardhat/deployments/sepolia/*.json` (ABI and compiler metadata per contract); `docs/data/` (measured Sepolia gas and events, the slice snapshot used by the reconciliation, local fixtures).
- **Financial analysis:** three Python notebooks in `docs/` (`analysis_pool.ipynb`, `analysis_slicing.ipynb`, `analysis_costs.ipynb`), each executed with its checks visible. The brief accepts a notebook in place of an Excel workbook.
- **Reproduce:** `yarn install`, `yarn test` (92 tests), `yarn demo-check` (read-only check of the Sepolia deployment); the run-book is `docs/sepolia-runbook.md`.

References

- Scaffold-ETH 2 documentation, https://scaffoldeth.io (project framework: Hardhat, Next.js, wagmi, viem, RainbowKit).
- OpenZeppelin Contracts v5 documentation, https://docs.openzeppelin.com/contracts (ERC-20, ERC-721, Ownable, ReentrancyGuard, SafeERC20).
- Chainlink Data Feeds documentation, https://docs.chain.link/data-feeds (the BTC/USD aggregator and its `latestRoundData` fields).
- Hardhat documentation, https://hardhat.org, and Slither, https://github.com/crytic/slither (testing and static analysis).

- Adams, Zinsmeister and Robinson, *Uniswap v2 Core*, 2020, https://uniswap.org/whitepaper.pdf (the constant-product rule, the 0.3% fee and the price accumulator that our pool and TWAP follow in spirit).
- Adams, White and Robinson, *Time-Weighted Automated Market Maker (TWAMM)*, Paradigm, 2021, https://www.paradigm.xyz/2021/07/twamm (related work; we deliberately use fixed slices instead of continuous execution).
- Impermanent loss: the closed form 2*sqrt(r)/(1+r) - 1 is standard; Section 4.4 checks it against the integer pool rather than relying on it.
- FRE-GY 6931 Final Project Brief, Fall 2026 (requirements, report headings, `submission.json` template).

The formulas and numbers in Sections 4 and 5 are derived and checked in this project's notebooks and tests; the notebooks do not cite external sources, so the list above is the technology we used and the background for the standard results.